
INFO 2950: Intro to Data Science

Lecture 14
2023-10-16

Agenda

1. Binary Evaluation Metrics

- a. Precision / Recall
- b. F1 / ROC-AUC

2. Train / test sets

- a. Cross Validation
- b. Sliding Windows

Precision and Recall

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- **Precision** = $tp / (tp + fp)$

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- **Recall** = $tp / (tp + fp)$

“True” Values Stay the Same

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

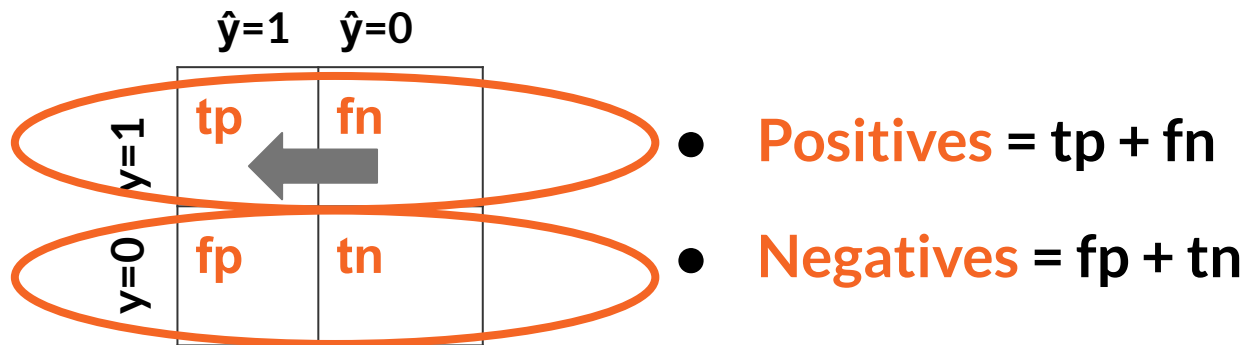
- **Positives** = tp + fn
- **Negatives** = fp + tn

“True” Values Stay the Same

	$\hat{y}=1$	$\hat{y}=0$	
$y=1$	tp	fn	• Positives = tp + fn
$y=0$	fp	tn	• Negatives = fp + tn

- If fn decreases, what happens to tp?

“True” Values Stay the Same



- If **fn** decreases, what happens to **tp**?
 - **tp increases since total positives must stay the same**

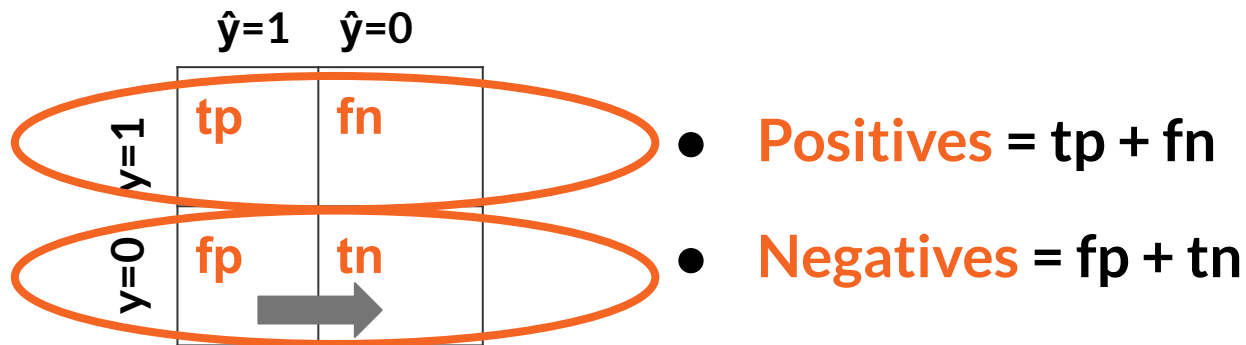
“True” Values Stay the Same

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

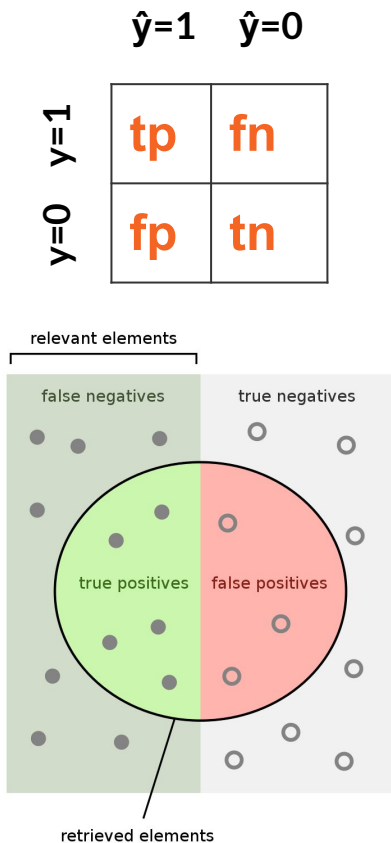
- **Positives** = tp + fn
- **Negatives** = fp + tn

- If tn increases, what happens to fp?

“True” Values Stay the Same



- If tn increases, what happens to fp?
 - **fp decreases since total negatives must stay the same**



How does recall change?

- Recall = $tp / (tp + fn)$
- If we change our model so the number of false negatives is decreased, what happens to recall?
 - Hint: let p represent the number of true $y=1$ values in your data

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- $\text{Recall} = \text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- Recall = $\text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

- Recall = $(p - \text{fn}) / (p)$

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- $\text{Recall} = \text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

- $\text{Recall} = (p - \text{fn}) / (p)$

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- $\text{Recall} = \text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

- $\text{Recall} = (p - \text{fn}) / (p)$

$$= 1 - (\text{fn}/p)$$

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- $\text{Recall} = \text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

- $\text{Recall} = (p - \text{fn}) / (p)$

$$= 1 - (\text{fn}/p)$$

- If fn decreases, (fn/p) decreases since p is constant

How does recall change?

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- $\text{Recall} = \text{tp} / (\text{tp} + \text{fn})$

- $p = \text{tp} + \text{fn}$

- $\text{Recall} = (p - \text{fn}) / (p)$

$$= 1 - (\text{fn}/p)$$

- If fn decreases, (fn/p) decreases since p is constant
- ...so, recall $[1 - (\text{fn}/p)]$ **always increases**

Before we talk about precision, let's talk about the model itself

	Model predicts 1	Model predicts 0
True value is 1	True positive Correct prediction	False negative
True value is 0	False positive	True negative Correct prediction

What model do we use for binary output?

	Model predicts 1	Model predicts 0
True value is 1	True positive Correct prediction	False negative
True value is 0	False positive	True negative Correct prediction

Logistic transformation: $\sigma(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$

General logistic function
where p denotes
probability that $y=1$:

$$p(x) = \sigma(t) = \frac{1}{1 + e^{-(\alpha + \beta \cdot x)}}$$

Logistic regression

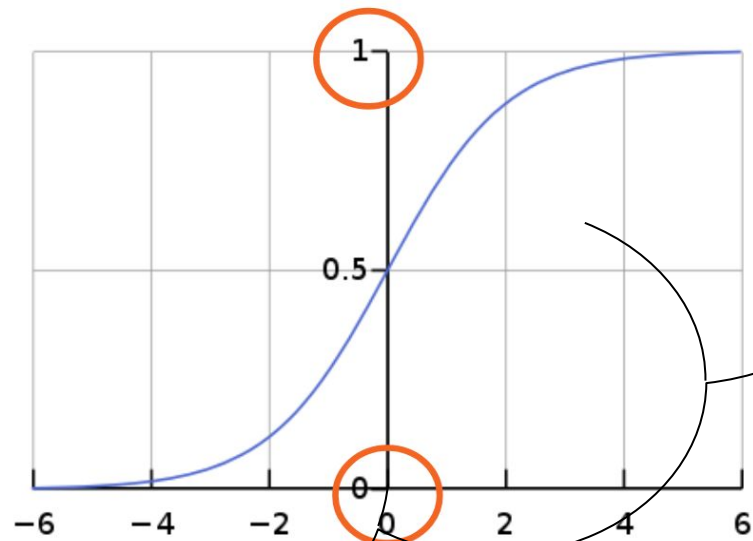
- When do we use it?
 - When output is

Logistic regression

- When do we use it?
 - When output is **binary**
- When you predict the \hat{y} for a given value of x , what do you get?
 - Is \hat{y} binary?
 - What's going on under the hood?

Logistic (sigmoid) function

All outputs of a logistic function could be a probability (between 0 and 1!)

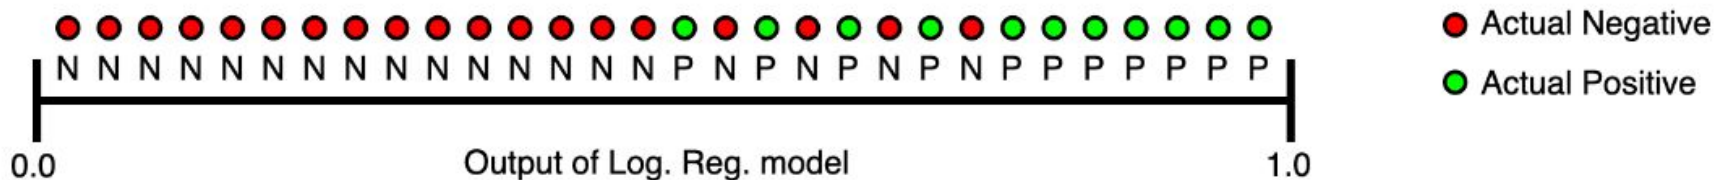


How do you get from probability to binary?

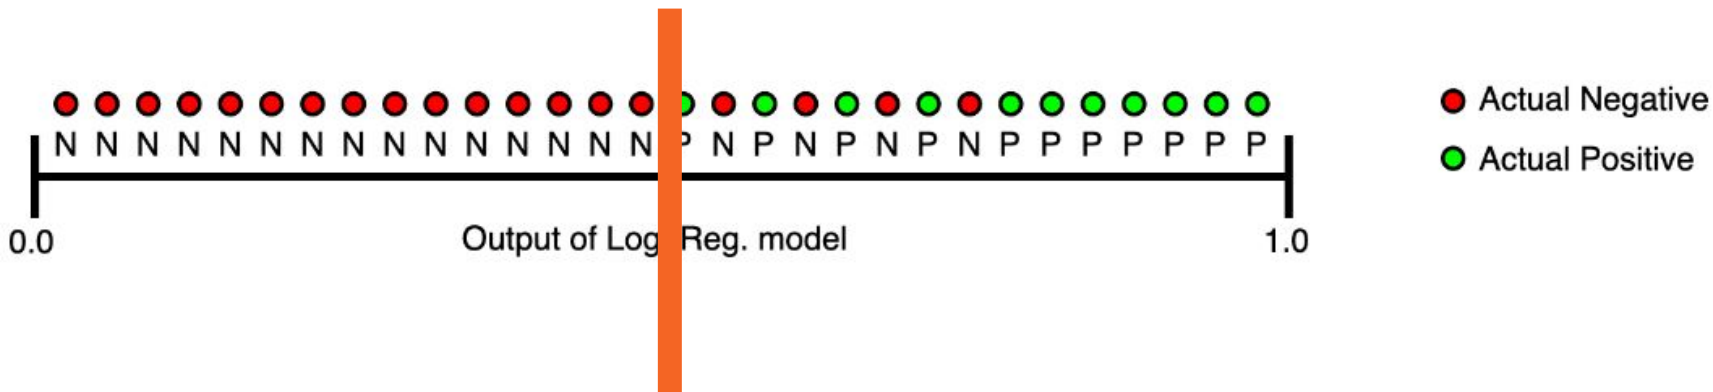
- Logistic regression interpretation (e.g. the “summarize” step) regards the odds of $y=1$
 - i.e., logit tells you about the **probability that $y=1$**

Sort the points by p (logistic output)

We have to make a decision.
How do we divide these?

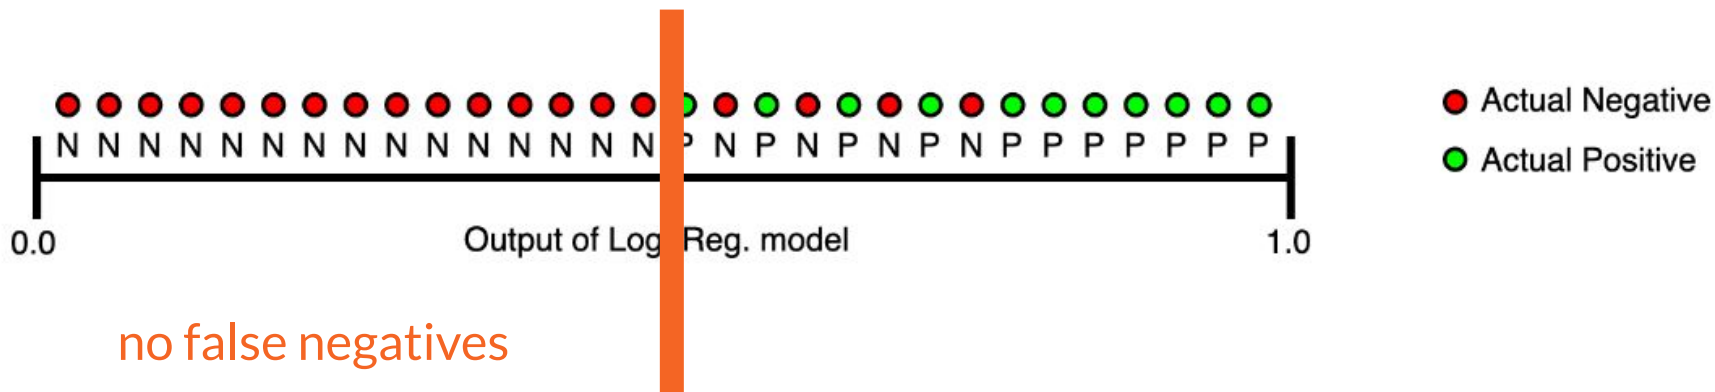


We have to make a decision.
How do we divide these?

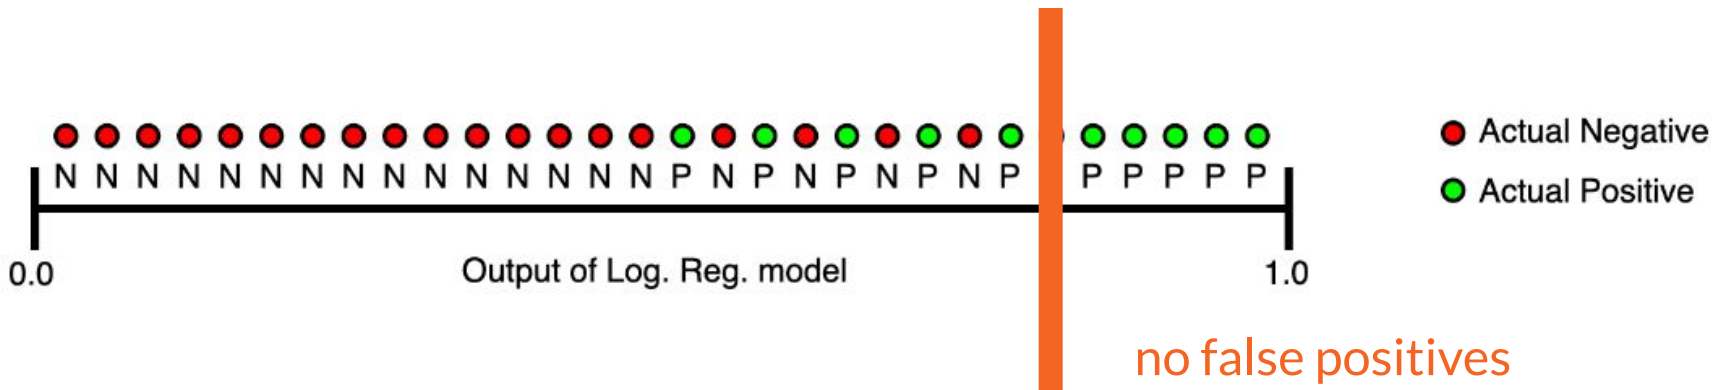


Sort the points by p (logistic output)

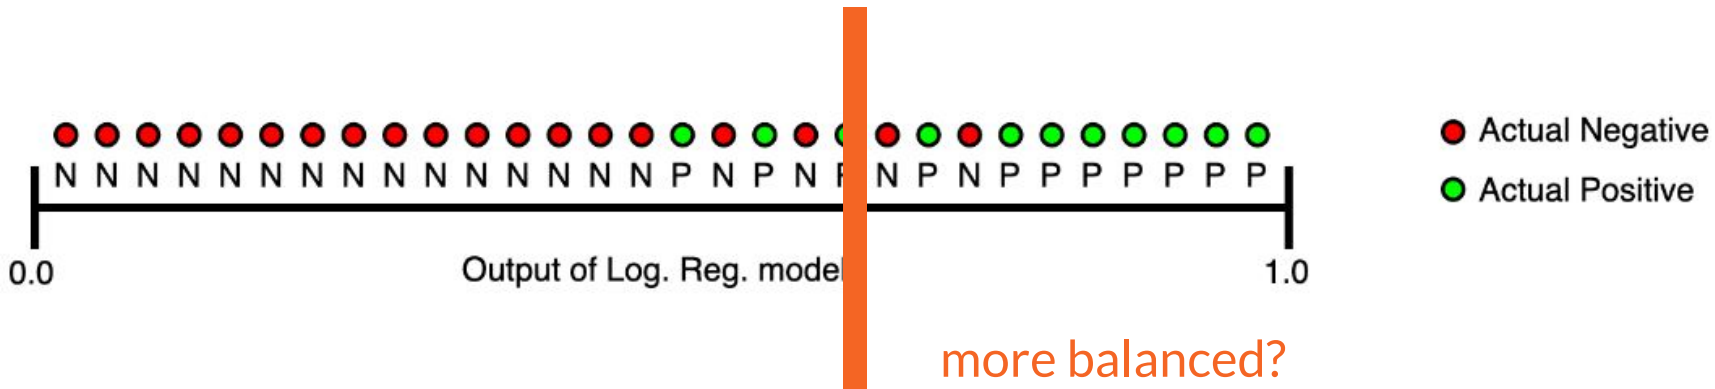
We have to make a decision.
How do we divide these?



We have to make a decision.
How do we divide these?



We have to make a decision.
How do we divide these?



What do we do with p ?

- You need to convert a **probability** (probability that $y=1$) to a **binary** output (y = either 1 or 0)
- If I tell you the probability of $y=1$ is 100%, what would you say that y = ?

What do we do with p ?

- You need to convert a **probability** (probability that $y=1$) to a **binary** output (y = either 1 or 0)
- If I tell you the probability of $y=1$ is 100%, what would you say that $y = ?$ (**1**)

What do we do with p ?

- You need to convert a **probability** (probability that $y=1$) to a **binary** output (y = either 1 or 0)
- If I tell you the probability of $y=1$ is **90%**, what would you say that y = ?

What do we do with p ?

- You need to convert a **probability** (probability that $y=1$) to a **binary** output (y = either 1 or 0)
- If I tell you the probability of $y=1$ is 90%, what would you say that $y = ?$ (**1**)

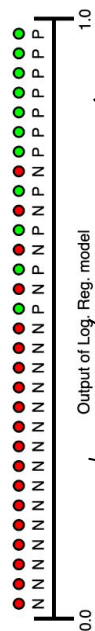
What do we do with p ?

- You need to convert a **probability** (probability that $y=1$) to a **binary** output (y = either 1 or 0)
- If I tell you the probability of $y=1$ is 90%, what would you say that y = ?
 - **What if the probability is 80, or 75, or 55, or 52, or 51, or 50.000001?**

Binarizing threshold

- Usually, you set a threshold of 50% to binarize a probability:
 - If you think there are better than 50/50 odds that $y=1$, then when choosing between outputting $y=1$ and $y=0$, you should choose $y=1$

● Actual Negative
● Actual Positive



Binarizing threshold

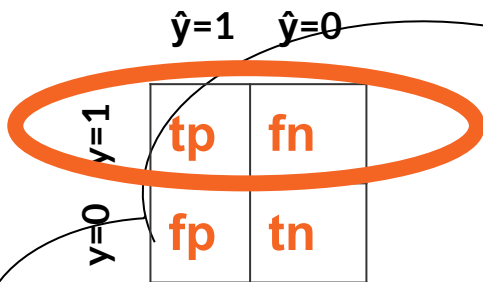
- What if we don't choose 50% as a threshold?

Image	\hat{y} Probability	50% Threshold	40% Threshold
Dog (1)	0.9	1	1
Dog (1)	0.4	0	1
Cat (0)	0.3	0	0

Changing thresholds

- Changing thresholds can let you change your precision and recall (possibly improving your metrics)
- We can use thresholds to show something about the relationship between precision and recall

Earlier: fn decrease \rightarrow recall increase



A confusion matrix diagram. The columns are labeled $\hat{y}=1$ and $\hat{y}=0$ at the top. The rows are labeled $y=1$ and $y=0$ on the left. The cells contain 'tp', 'fn', 'fp', and 'tn' respectively. An orange oval highlights the top row (tp and fn). A curved line connects the 'tp' cell to the first bullet point. Another curved line connects the 'fn' cell to the second and third bullet points. A third curved line connects the 'fp' cell to the fourth bullet point. There are also some small circles at the bottom left of the slide.

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

- Recall = $tp / (tp + fn)$
- $p = tp + fn$
- Recall = $(p - fn) / (p)$

$$= 1 - (fn/p)$$

- If fn decreases, (fn/p) decreases since p is constant
- ...so, recall $[1 - (fn/p)]$ **always increases**

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

fn decrease \rightarrow precision change?

- Precision = $tp / (tp + fp)$
- If we change thresholds and decrease the number of false negatives produced by our model, what happens to precision?

y (0 or 1)	\hat{y} (probability)	Threshold 1	Threshold 2
...

Decreasing fn can decrease precision = $tp/(tp+fp)$

Image	\hat{y} Probability	50% Threshold
Dog (1)	0.8	1
Dog (1)	0.2	0
Cat (0)	0.2	0
Cat (0)	0.2	0

Decreasing fn can decrease precision = $tp/(tp+fp)$

Image	\hat{y} Probability	50% Threshold (1 false negative)
Dog (1)	0.8	1
Dog (1)	0.2	0
Cat (0)	0.2	0
Cat (0)	0.2	0

Decreasing fn can decrease precision = $tp / (tp + fp)$

$$\text{Precision} = 1 / (1 + 0) = 1$$

Image	\hat{y} Probability	50% Threshold (1 false negative)
Dog (1)	0.8	1
Dog (1)	0.2	0
Cat (0)	0.2	0
Cat (0)	0.2	0

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

There are 0 false positives (cats predicted as dogs)

Decreasing fn can decrease precision = $tp/(tp+fp)$

$$\text{Precision} = 1 / (1+0) = 1$$

Image	\hat{y} Probability	50% Threshold (1 false negative)	0% Threshold (0 false negatives)
Dog (1)	0.8	1	?
Dog (1)	0.2	0	?
Cat (0)	0.2	0	?
Cat (0)	0.2	0	?

$$\text{Precision} = ?$$

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

Decreasing fn can decrease precision = $tp / (tp + fp)$

$$\text{Precision} = 1 / (1 + 0) = 1$$

Image	\hat{y} Probability	50% Threshold (1 false negative)	0% Threshold (0 false negatives)
Dog (1)	0.8	1	1
Dog (1)	0.2	0	1
Cat (0)	0.2	0	1
Cat (0)	0.2	0	1

Any probability > threshold = 1
All probabilities > 0 → Everything = 1!

Decreasing fn can decrease precision = $tp / (tp + fp)$

$$\text{Precision} = 1 / (1 + 0) = 1$$

Image	\hat{y} Probability	50% Threshold (1 false negative)	0% Threshold (0 false negatives)
Dog (1)	0.8	1	1
Dog (1)	0.2	0	1
Cat (0)	0.2	0	1
Cat (0)	0.2	0	1

$$\text{Precision} = 2 / (2 + 2) = 0.5$$

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

Decreasing fn can decrease precision = $tp/(tp+fp)$

Precision = $1 / (1+0) = 1$ \longrightarrow Precision = $2 / (2+2) = 0.5$

Image	\hat{y} Probability	50% Threshold (1 false negative) \longrightarrow	0% Threshold (0 false negatives)
Dog (1)	0.8	1	1
Dog (1)	0.2	0	1
Cat (0)	0.2	0	1
Cat (0)	0.2	0	1

But decreasing fn can also **increase** precision = $tp / (tp + fp)$

$$\text{Precision} = 0 / (0 + 1) = 0$$

Image	\hat{y} Probability	50% Threshold (1 false negative)	0% Threshold (0 false negatives)
Cat (0)	1.0	1	?
Dog (1)	0.4	0	?
Cat (0)	0.1	0	?

$$\text{Precision} = ?$$

But decreasing fn can also
increase precision = $tp / (tp + fp)$

$$\text{Precision} = 0 / (0 + 1) = 0$$

Image	\hat{y} Probability	50% Threshold (1 false negative)	0% Threshold (0 false negatives)
Cat (0)	1.0	1	1
Dog (1)	0.4	0	1
Cat (0)	0.1	0	1

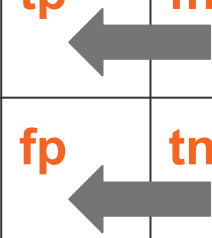
$$\text{Precision} = 1 / (1 + 2) = 1/3$$

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn

Changing thresholds

Decreasing the threshold (e.g. from 50% to 0%) can shift your false negatives and your true negatives towards positives

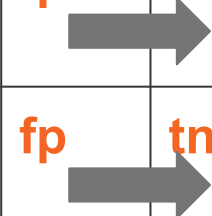
	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn



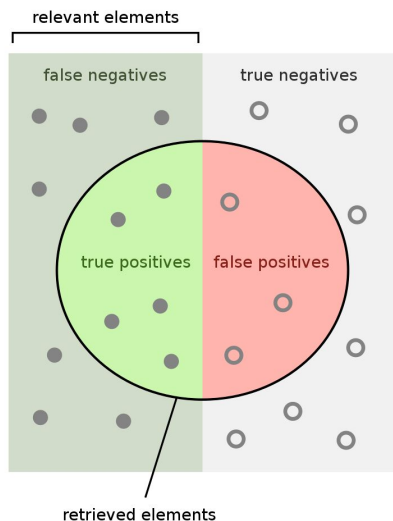
Changing thresholds

Increasing the threshold (e.g. from 0% to 50%) can shift your false positives and your true positives towards negatives

	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn



	$\hat{y}=1$	$\hat{y}=0$
$y=1$	tp	fn
$y=0$	fp	tn



Precision and Recall Recap

- If we change the threshold to decrease the number of false negatives ...
 - Recall necessarily increases
 - Precision could increase or decrease

Admin

- Phase 2 due this Thursday at 11:59 p.m.
 - No slip days are allowed for project deadlines
- Prelim regrade requests must be made in Gradescope by the end of October
- If project groups need to change after drop deadline, post on Ed to find a group / new members

A classification model separates email into two categories: "spam" or "not spam." If you raise the classification threshold, what will happen to precision? [Precision = $tp / (tp+fp)$]

1. Probably Decrease
2. Definitely Decrease
3. Probably Increase
4. Definitely Increase

A classification model separates email into two categories: "spam" or "not spam." If you raise the classification threshold, what will happen to precision?

1. Probably Decrease
2. Definitely Decrease
3. Probably Increase
4. Definitely Increase

"In general, raising the classification threshold reduces false positives, thus raising precision."

A classification model separates email into two categories: "spam" or "not spam." If you raise the classification threshold, what will happen to recall? [Recall = $tp / (tp + fn)$]

1. Always increase
2. Always stay constant
3. Always decrease or stay the same

A classification model separates email into two categories: "spam" or "not spam." If you raise the classification threshold, what will happen to recall?

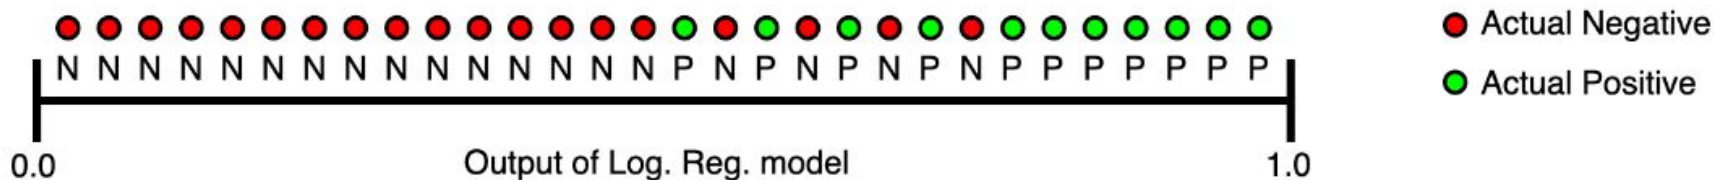
1. Always increase
2. Always stay constant
3. Always decrease or stay the same

“Raising our classification threshold will cause the number of true positives to decrease or stay the same and will cause the number of false negatives to increase or stay the same. Thus, recall will either stay constant or decrease.”

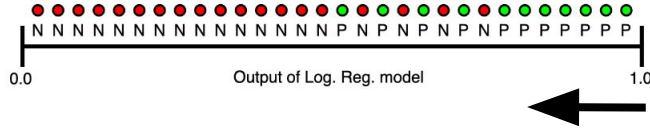


- What can we say about the relationship between recall and precision, both of which we want to be high?

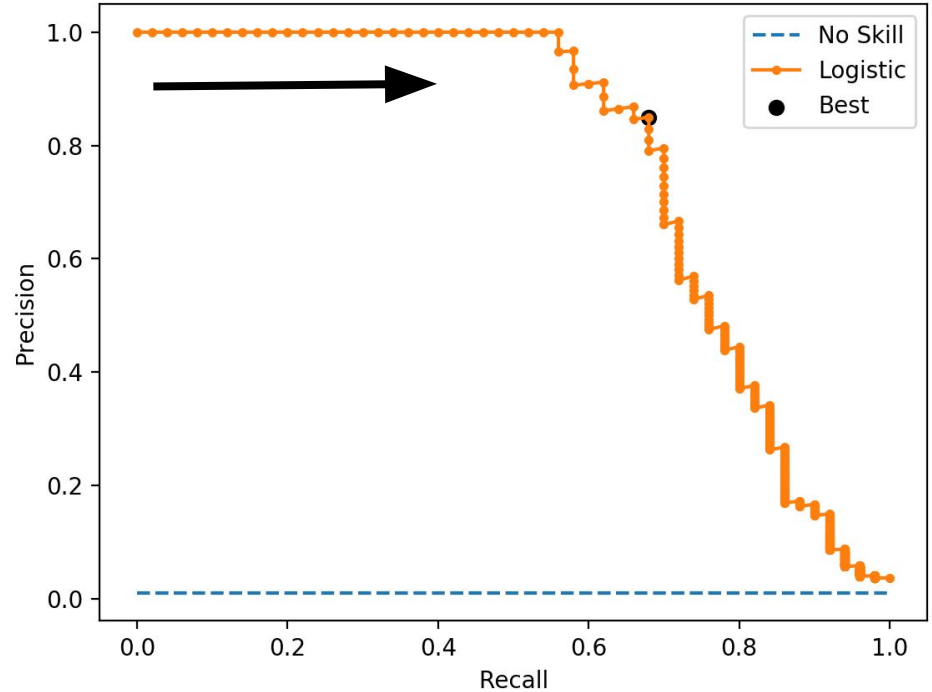
What if we consider a lot of thresholds all at once?



Precision Recall Curves

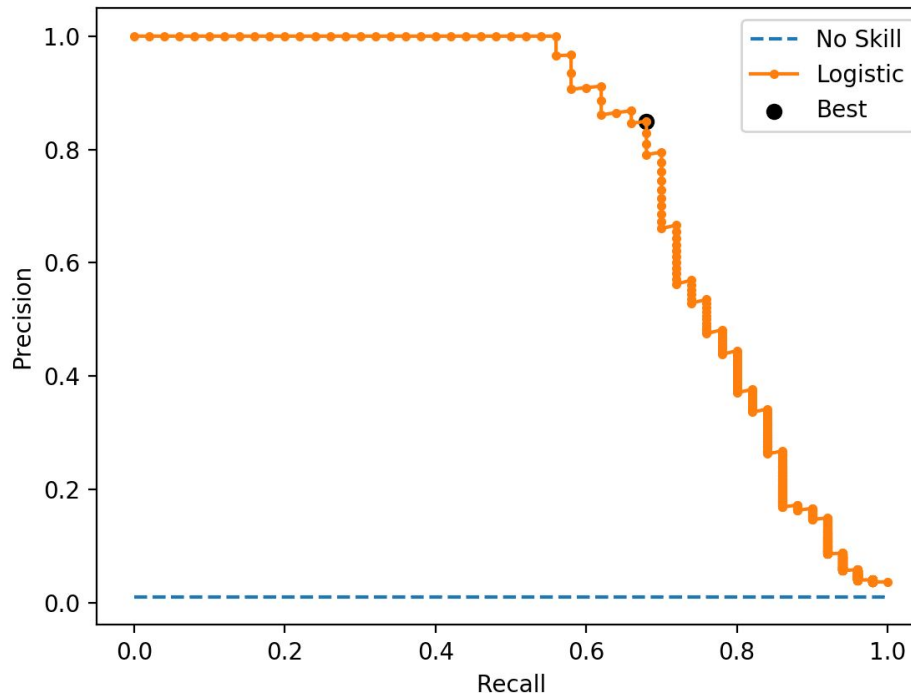


● Actual Negative
● Actual Positive



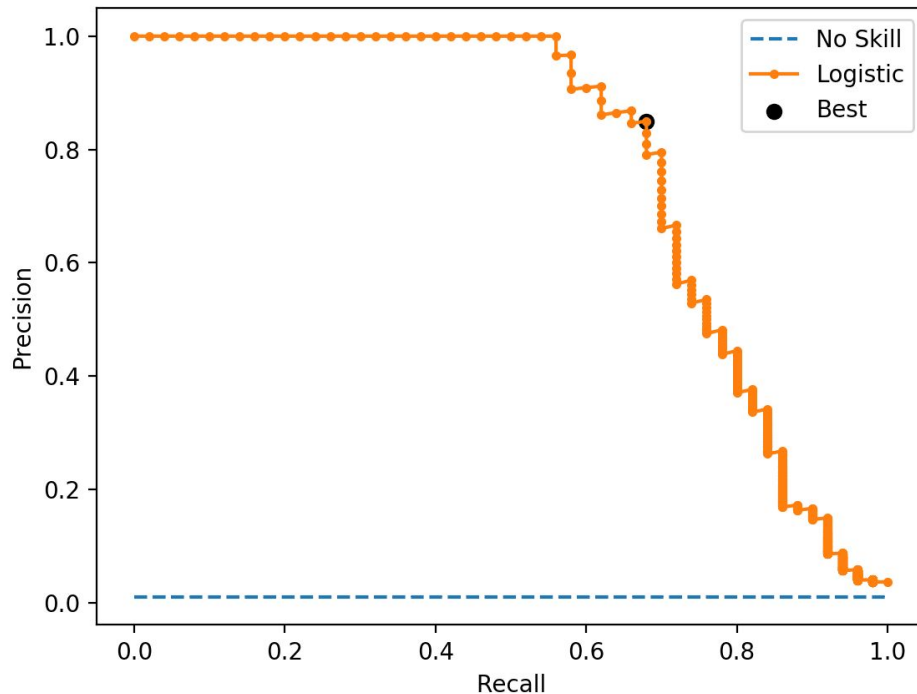
Precision Recall Curves

In general, for a model better than random, recall and precision have an inverse relationship



Precision Recall Curves

You can plot PR curves to help you think about what threshold is “Best” (giving you ideal achievable recall and precision)



F1 combines precision and recall

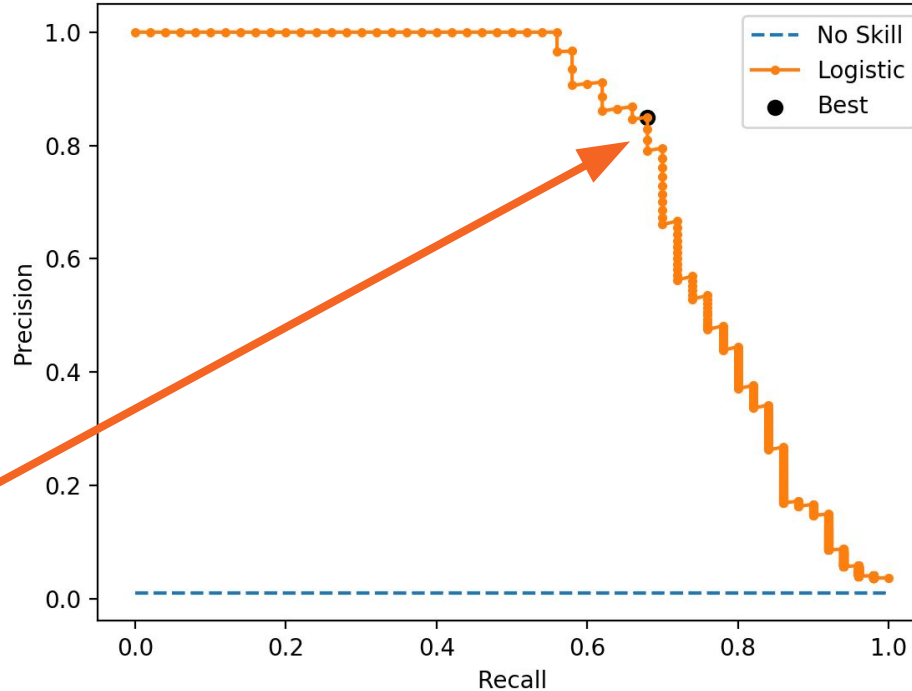
- **F₁ score:** $\frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}} = \frac{2tp}{2tp + fp + fn}$
- This is the *harmonic mean* of precision and recall
- Higher F₁ is better

(Lots of options for classification)

Sources: [24][25][26][27][28][29][30][31][32] view · talk · edit

		Predicted condition			
		Positive (PP)	Negative (PN)		
Actual condition	Total population = P + N			Informedness, bookmaker informedness (BM) = TPR + TNR - 1	Prevalence threshold (PT) $= \frac{\sqrt{TPR \times FPR} - FPR}{TPR - FPR}$
	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power = $\frac{TP}{P} = 1 - FNR$	False negative rate (FNR), miss rate = $\frac{FN}{P} = 1 - TPR$
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection	False positive rate (FPR), probability of false alarm, fall-out = $\frac{FP}{N} = 1 - TNR$	True negative rate (TNR), specificity (SPC), selectivity = $\frac{TN}{N} = 1 - FPR$
	Prevalence = $\frac{P}{P + N}$	Positive predictive value (PPV), precision = $\frac{TP}{PP} = 1 - FDR$	False omission rate (FOR) = $\frac{FN}{PN} = 1 - NPV$	Positive likelihood ratio (LR+) = $\frac{TPR}{FPR}$	Negative likelihood ratio (LR-) = $\frac{FNR}{TNR}$
	Accuracy (ACC) = $\frac{TP + TN}{P + N}$	False discovery rate (FDR) = $\frac{FP}{PP} = 1 - PPV$	Negative predictive value (NPV) = $\frac{TN}{PN} = 1 - FOR$	Markedness (MK), deltaP (Δp) = PPV + NPV - 1	Diagnostic odds ratio (DOR) = $\frac{LR+}{LR-}$
	Balanced accuracy (BA) = $\frac{TPR + TNR}{2}$	F₁ score = $\frac{2PPV \times TPR}{PPV + TPR} = \frac{2TP}{2TP + FP + FN}$	Fowlkes–Mallows index (FM) = $\sqrt{PPV \times TPR}$	Matthews correlation coefficient (MCC) = $\sqrt{TPR \times TNR \times PPV \times NPV} - \sqrt{FNR \times FPR \times FOR \times FDR}$	Threat score (TS), critical success index (CSI), Jaccard index = $\frac{TP}{TP + FN + FP}$

F1 score



This is the point that maximizes the F_1 score (harmonic mean of precision and recall)

Binary evaluation metrics, part 2

- **F₁ score:** $\frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}} = \frac{2\text{tp}}{2\text{tp} + \text{fp} + \text{fn}}$
- **AUC-ROC curve:** fp / (tn + fp) a.k.a. specificity versus recall
- Higher F₁ and AUC is better



Another curve: ROC

- **Receiver Operating Characteristic (ROC)** plots the FPR vs. TPR
 - Originally developed for military radar purposes



What is ROC?

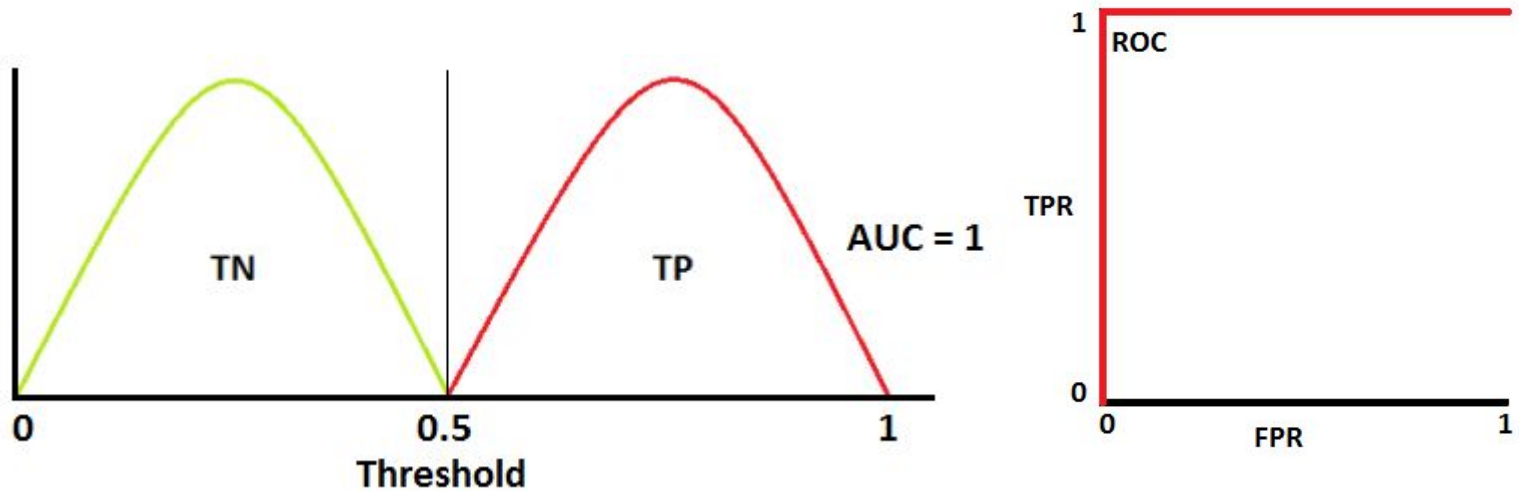
- **Receiver Operating Characteristic (ROC)** plots the FPR vs. TPR
 - Originally developed for military radar purposes
- **False Positive Rate** (a.k.a. 1-specificity)
 - $fp / (fp + tn)$
- **True Positive Rate** (a.k.a. recall, sensitivity)
 - $tp / (tp + fn)$

(Lots of options for classification)

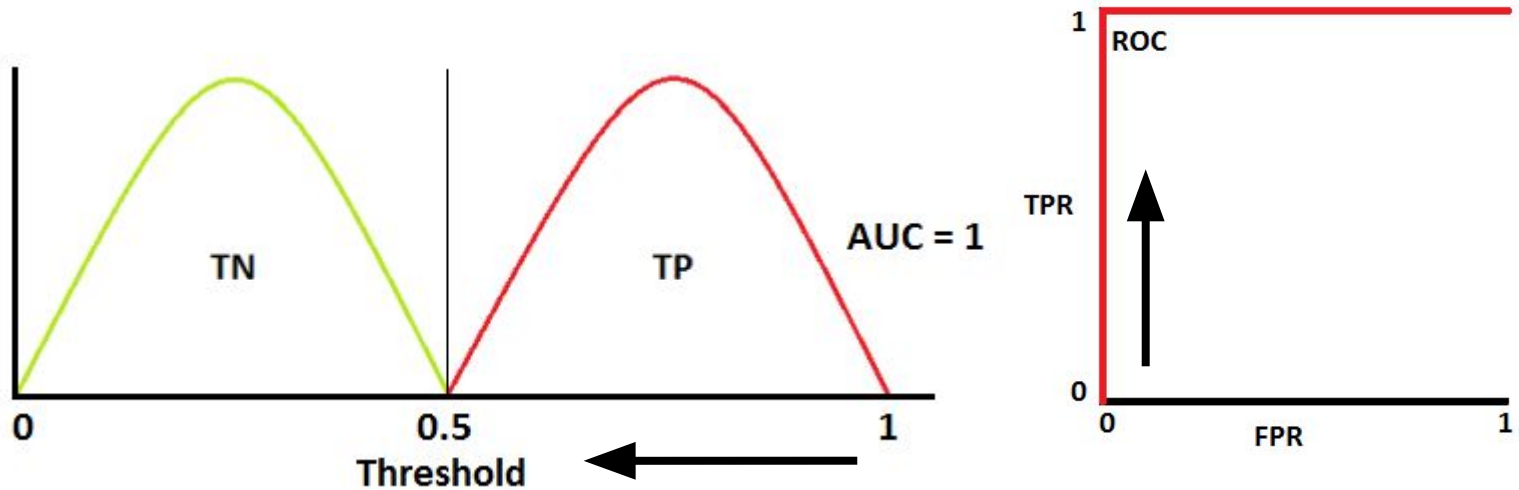
Sources: [24][25][26][27][28][29][30][31][32] view · talk · edit

		Predicted condition			
		Positive (PP)	Negative (PN)	Informedness, bookmaker informedness (BM) $= \text{TPR} + \text{TNR} - 1$	Prevalence threshold (PT) $= \frac{\sqrt{\text{TPR} \times \text{FPR}} - \text{FPR}}{\text{TPR} - \text{FPR}}$
Actual condition	Total population $= P + N$				
	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power $= \frac{\text{TP}}{P} = 1 - \text{FNR}$	False negative rate (FNR), miss rate $= \frac{\text{FN}}{P} = 1 - \text{TPR}$
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection	False positive rate (FPR), probability of false alarm, fall-out $= \frac{\text{FP}}{N} = 1 - \text{TNR}$	True negative rate (TNR), specificity (SPC), selectivity $= \frac{\text{TN}}{N} = 1 - \text{FPR}$
		Positive predictive value (PPV), precision $= \frac{\text{TP}}{\text{PP}} = 1 - \text{FDR}$	False omission rate (FOR) $= \frac{\text{FN}}{\text{PN}} = 1 - \text{NPV}$	Positive likelihood ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$	Negative likelihood ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$
		Accuracy (ACC) $= \frac{\text{TP} + \text{TN}}{P + N}$	False discovery rate (FDR) $= \frac{\text{FP}}{\text{PP}} = 1 - \text{PPV}$	Negative predictive value (NPV) $= \frac{\text{TN}}{\text{PN}} = 1 - \text{FOR}$	Markedness (MK), deltaP (Δp) $= \text{PPV} + \text{NPV} - 1$
		Balanced accuracy (BA) $= \frac{\text{TPR} + \text{TNR}}{2}$	F_1 score $= \frac{2\text{PPV} \times \text{TPR}}{\text{PPV} + \text{TPR}} = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}}$	Fowlkes–Mallows index (FM) $= \sqrt{\text{PPV} \times \text{TPR}}$	Matthews correlation coefficient (MCC) $= \sqrt{\text{TPR} \times \text{TNR} \times \text{PPV} \times \text{NPV}} - \sqrt{\text{FNR} \times \text{FPR} \times \text{FOR} \times \text{FDR}}$
					Threat score (TS), critical success index (CSI), Jaccard index $= \frac{\text{TP}}{\text{TP} + \text{FN} + \text{FP}}$

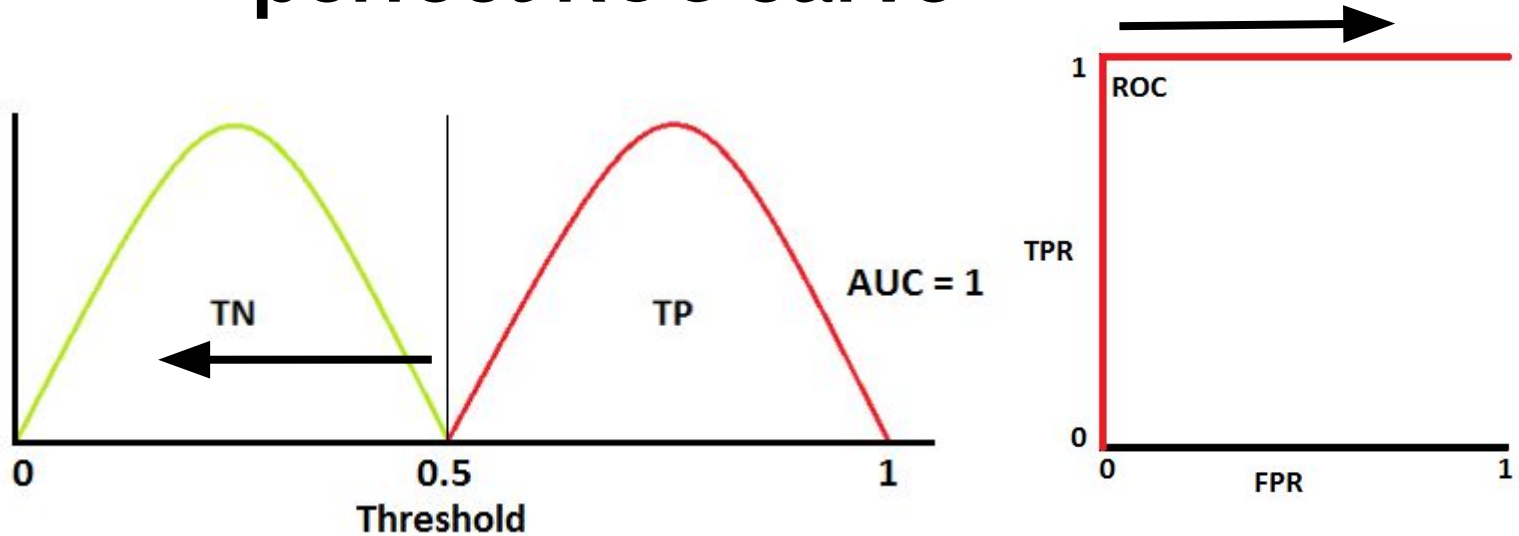
Perfectly separated data → perfect ROC curve



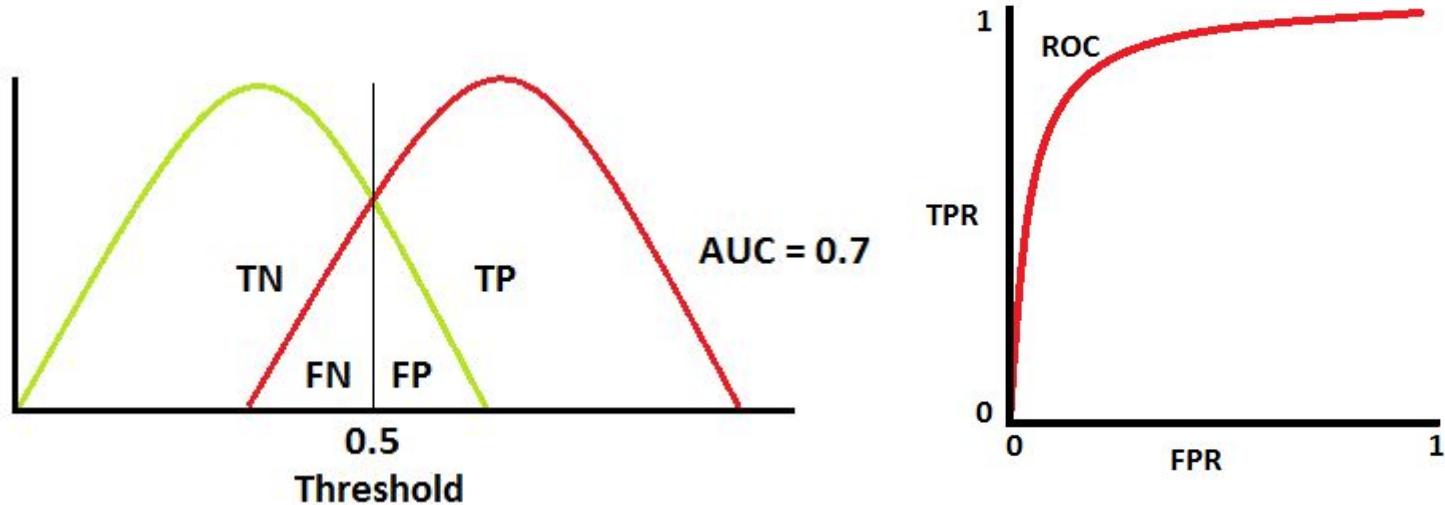
Perfectly separated data → perfect ROC curve



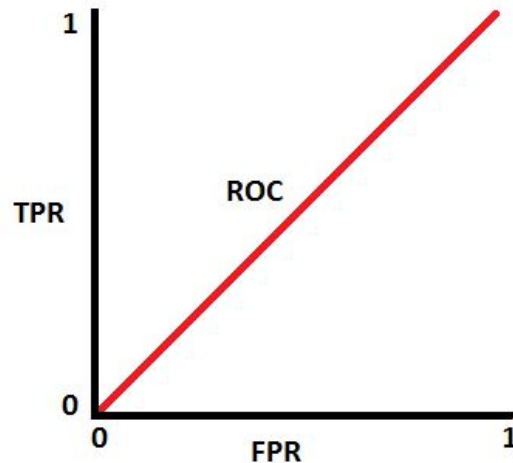
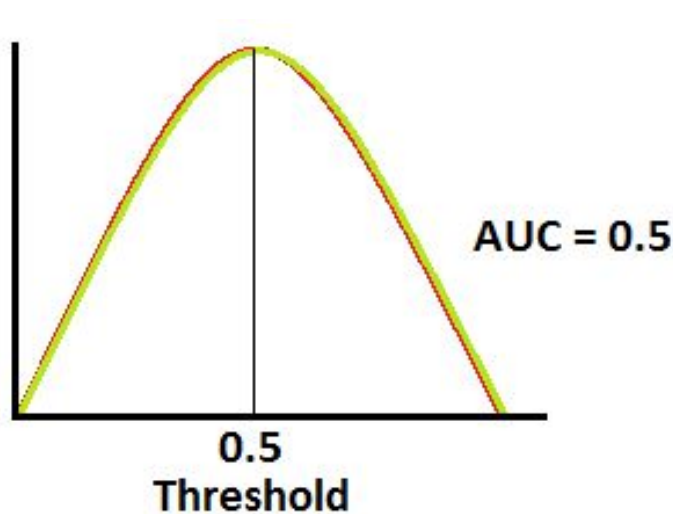
Perfectly separated data → perfect ROC curve



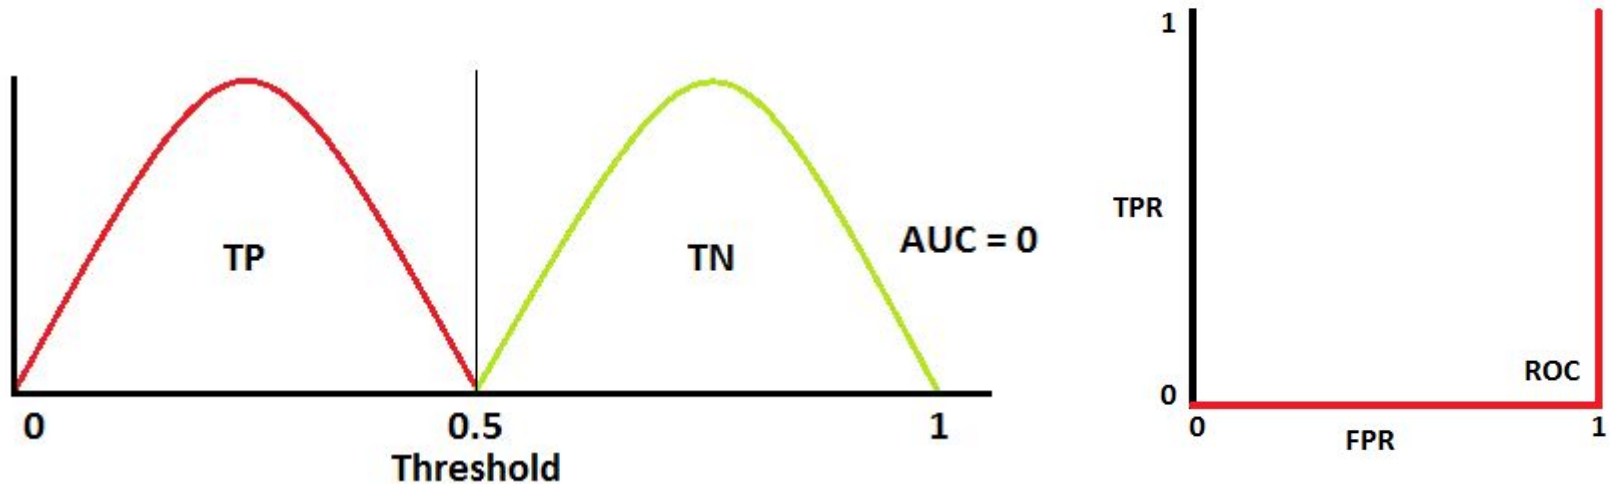
**Slightly overlapping data →
slightly lower ROC curve**



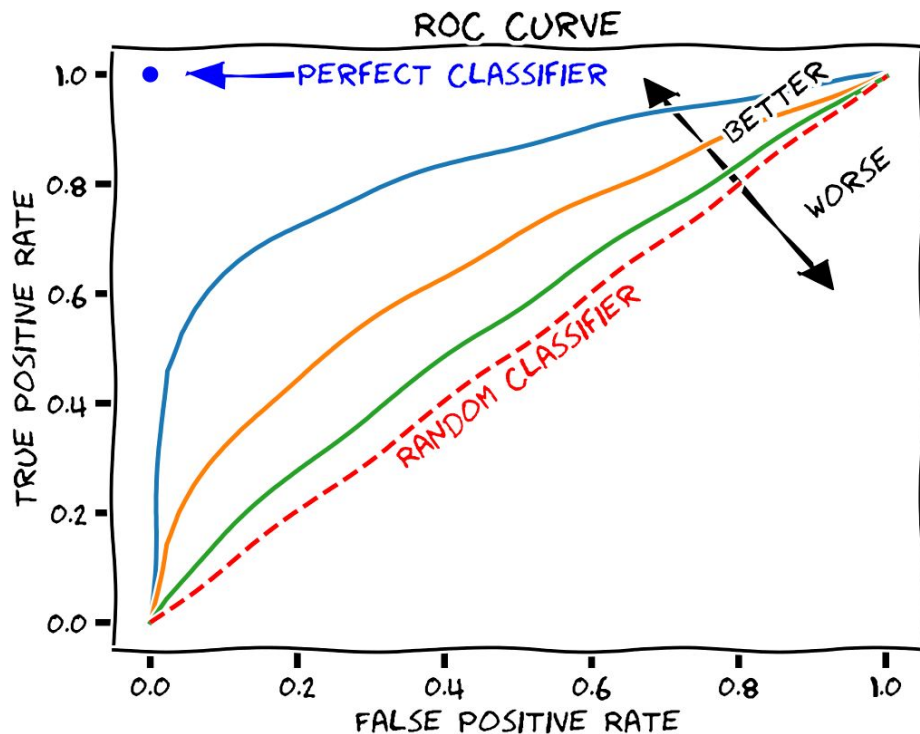
Exactly overlapping data → diagonal ROC curve (as good as random)



Perfectly flipped data → worst ROC curve



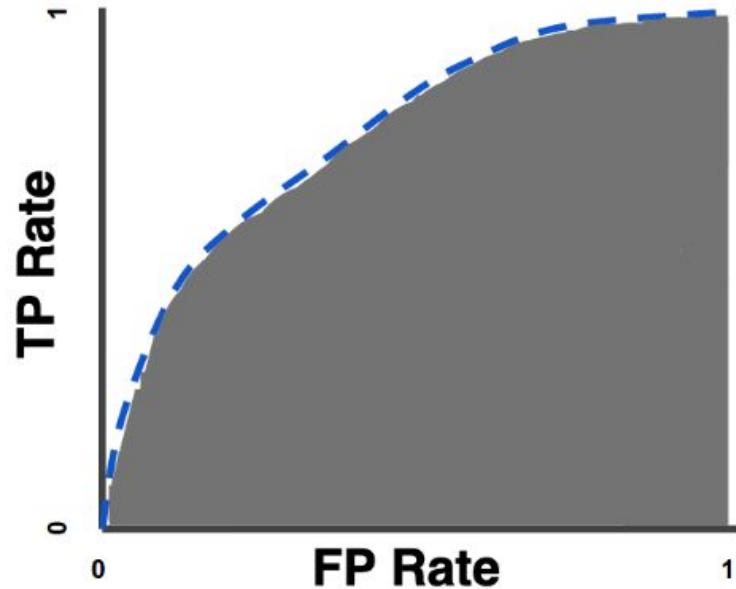
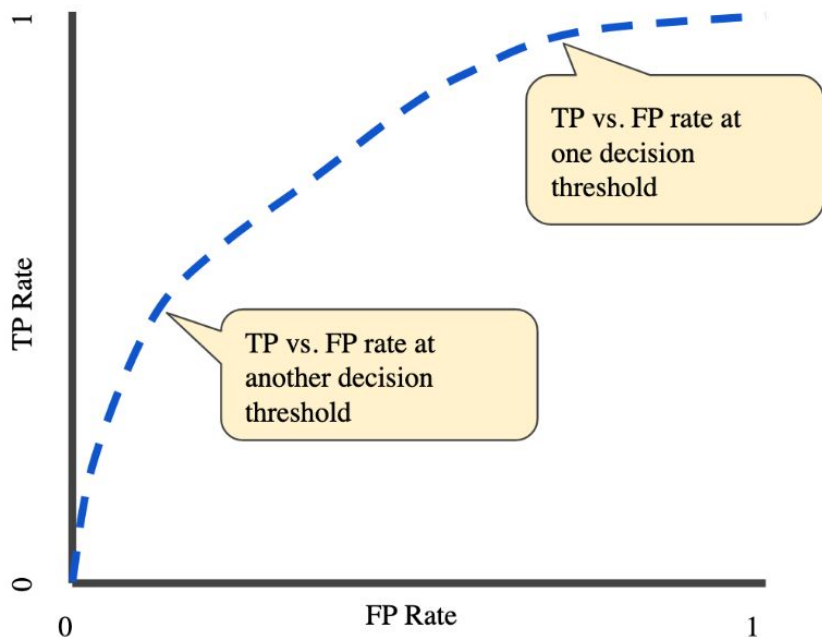
Good vs. Bad ROC



ROC curve



AUC



AUC (Area under the ROC Curve).

- Ability of your classifier to distinguish between two classes
- Summarizes the ROC curve



Probability that your logit model ranks a random positive example more highly than a random negative example



Classification prediction metrics in Python

- `from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score`
- `precision_score(y_true, y_hat)`
- `recall_score(y_true, y_hat)`
- `f1_score(y_true, y_hat)`
- `roc_auc_score(y_true, y_hat)`

Nomenclature for these slides

When providing
formulas in these
slides, we'll
generalize to these

General case	"y_true", "y_i"	"y_hat", "\hat{y}_i"
Train set	y_train	y_hat_train
Test set	y_test	y_hat_test

Classification prediction metrics in Python

- `from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score`
- `precision_score(y_true, y_hat)`
- `recall_score(y_true, y_hat)`
- `f1_score(y_true, y_hat)`
- `roc_auc_score(y_true, y_hat)`

Check the Scikit Learn documentation!



Install User Guide API Examples Community More ▾

 Go

Prev Up Next

scikit-learn 1.1.2

[Other versions](#)

Please [cite us](#) if you use the software.

3.3. Metrics and scoring: quantifying the quality of predictions

3.3.1. The **scoring** parameter: defining model evaluation rules

3.3.2. Classification metrics

3.3.3. Multilabel ranking metrics

3.3.4. Regression metrics

3.3.5. Clustering metrics

3.3.6. Dummy estimators

3.3. Metrics and scoring: quantifying the quality of predictions

There are 3 different APIs for evaluating the quality of a model's predictions:

- **Estimator score method:** Estimators have a `score` method providing a default evaluation criterion for the problem they are designed to solve. This is not discussed on this page, but in each estimator's documentation.
- **Scoring parameter:** Model-evaluation tools using [cross-validation](#) (such as `model_selection.cross_val_score` and `model_selection.GridSearchCV`) rely on an internal *scoring* strategy. This is discussed in the section [The scoring parameter: defining model evaluation rules](#).
- **Metric functions:** The `sklearn.metrics` module implements functions assessing prediction error for specific purposes. These metrics are detailed in sections on [Classification metrics](#), [Multilabel ranking metrics](#), [Regression metrics](#) and [Clustering metrics](#).

Finally, [Dummy estimators](#) are useful to get a baseline value of those metrics for random predictions.

See also: For "pairwise" metrics, between *samples* and not estimators or predictions, see the [Pairwise metrics, Affinities and Kernels](#) section.

1 minute break & attendance



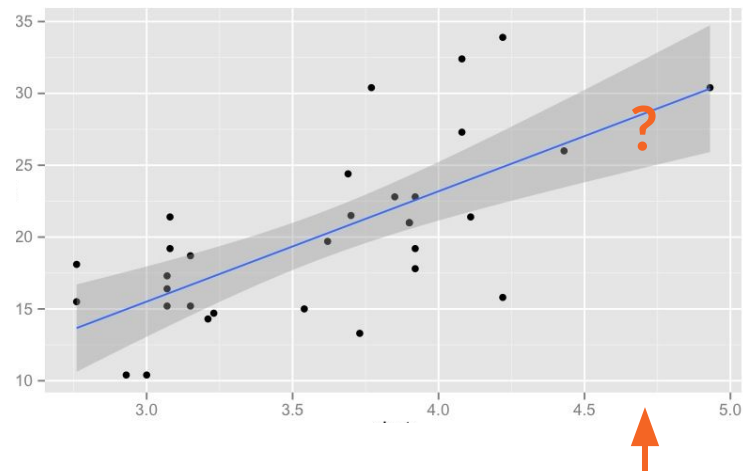
tinyurl.com/ykv3f9a9



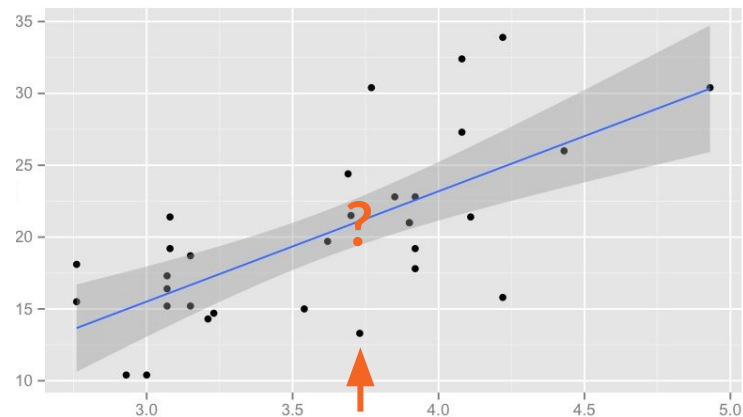
<https://pbs.twimg.com/media/DdkUUTMV4AAsohU.jpg>

Motivation

- How confident can we be that the patterns in the things we *have* seen will apply to the things we *haven't* seen?



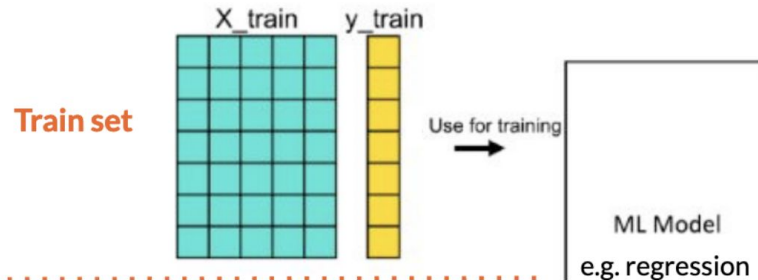
Motivation



- How confident can we be that the patterns in the things we *have* seen will apply to the things we *haven't* seen?
- **Key insight:** pretend we *haven't* seen some of the things we *have* seen, and check if we would have guessed correctly

What do you do with train / test sets?

- **Step 1:** experiment with your regression on your training set. Make any adjustments you need to here (e.g. try different models, transformations, etc.)



```
model = LinearRegression().fit(X_train,y_train)
```

<https://builtin.com/data-science/train-test-split> ⁸⁵

What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)

What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)

What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)

What could go wrong with train/test?

If your train set is only side sleepers and your test set is only back sleepers, no matter how you change your model, you still won't get good evaluations and will overfit to side sleepers!

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)

What could go wrong with train/test?

If your original set is only US undergrads, no amount of splitting will test generalization to middle-aged Brazilian manufacturing workers

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)

What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)
- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

What could go wrong with train/test?

If you split your time series data randomly into train and test data, you'll have time series with a bunch of random missing days. Missing chronological data is bad!

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)
- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

What could go wrong with train/test?

Solution:
Cross-Validation

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)
- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)

Solution:
Sliding Windows

- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

StackExchange

Search on Cross Validated...

Cross Validated is a question and answer site for people interested in statistics, machine learning, data analysis, data mining, and data visualization. It only takes a minute to sign up.

Sign up to join this community

Anybody can ask a question

Anybody can answer

The best answers are voted up and rise to the top

CrossValidated

Home

PUBLIC

Questions

Tags

Users

Unanswered

TEAMS

Stack Overflow for Teams – Start

Explore our questions

Ask Question

r

regression

machine-learning

time-series

probability

hypothesis-testing

distributions

self-study

neural-networks

bayesian

more tags

Active

4

Bountied

Hot

Week

Month

3 votes

3 answers

r

central-limit-theorem

405 views

Central-limit theorem via sample size or sampling magnitude?

Community Bot

1 modified 1 min ago

95

StackExchange

Search on Cross Validated...


Cross Validated is a question and answer site for people interested in statistics, machine learning, data analysis, data mining, and data visualization. It only takes a minute to sign up.

Sign up to join this community

Anybody can ask a question

Anybody can answer

The best answers are voted up and rise to the top

 Cross Validated

Home

PUBLIC

Questions

Tags

Users

Unanswered

TEAMS

Stack Overflow for Teams – Start

Explore our questions

Ask Question

r

regressionmachine-learningtime-seriesprobabilityhypothesis-testingdistributionsself-studyneural-networksbayesianmore tags

Active4BountiedHotWeekMonth

3 votes

3 answers

r

central-limit-theorem

405 views

Central-limit theorem via sample size or sampling magnitude?

CommunityBot1 modified 1 min ago

Validation Sets

- We know we can split data into train / test sets
- But there are a lot of reasons this may not be reliable
- What if we do the train / test splitting process *again* with the training set?

We have lots of choices of prediction models

- Which variables should we include as inputs?
- Should we use transformations?
- Should we filter training data?
- Should we use more complicated models? Interactions?

We have lots of choices of prediction models

- Which variables should we include as inputs?
- Should we use transformations?
- Should we filter training data?
- Should we use more complicated models? Interactions?

Danger: The more models we try, the more chances to get good results by pure luck

**The more tries
you take, the
more likely
you'll get good
eval metrics**

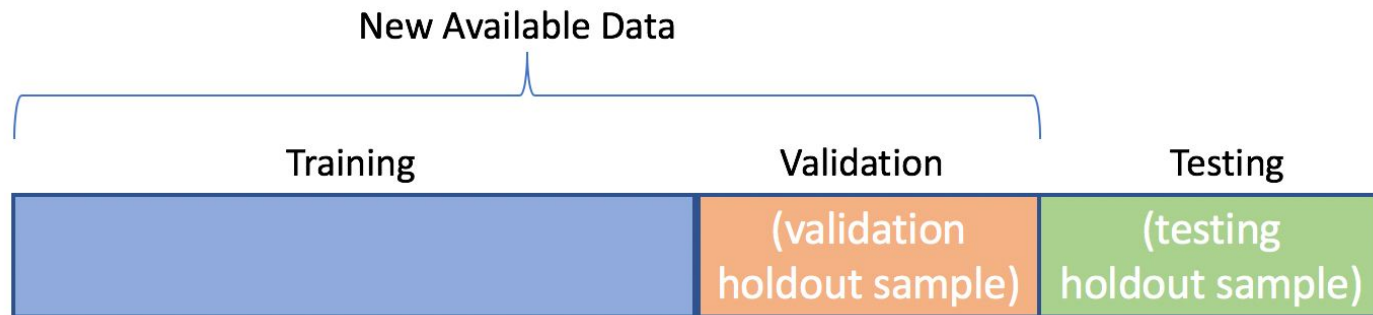


Validation Sets

- Don't want to cheat by overusing your test set
- You can only use your test set so many times before it's no longer useful to discover overfitting (since you'll start fitting to your test set too)
- Use validation set to check your evaluation metrics before checking your “final” test set

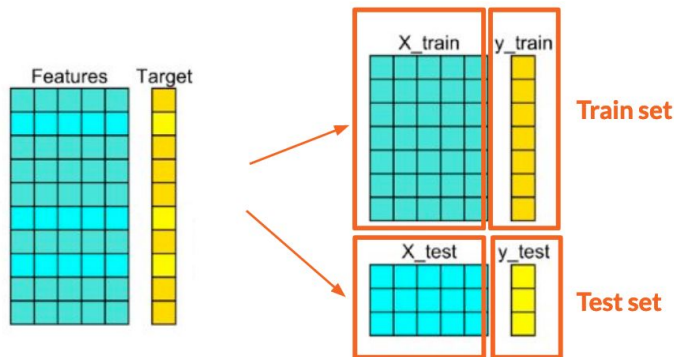
Available Data





Sub-split training to get validation

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3)
```



Sub-split training to get validation

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test =
```

```
    train_test_split(X, y, test_size = 0.3)
```

```
X_train, X_valid, y_train, y_valid =
```

```
    train_test_split(X_train, y_train, test_size = 0.1)
```



Sub-split training to get validation

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3)  
X_train, X_valid, y_train, y_valid =  
    train_test_split(X_train, y_train, test_size = 0.1)
```

Sub-split training to get validation

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3) } 70/30 train/test  
                                           } split in first step  
X_train, X_valid, y_train, y_valid =  
    train_test_split(X_train, y_train, test_size = 0.1) } ?
```

What % of the overall dataset are your train / valid / test sets?

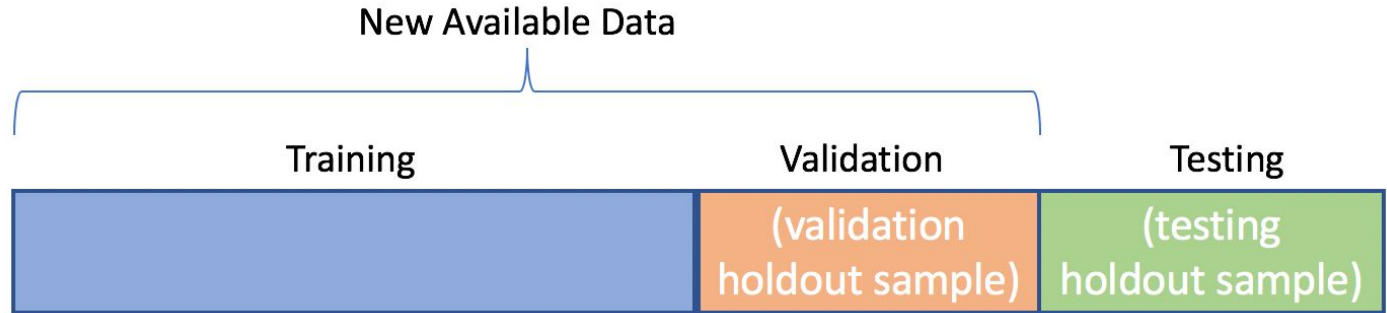
Sub-split training to get validation

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3)  
  
X_train, X_valid, y_train, y_valid =  
    train_test_split(X_train, y_train, test_size = 0.1)
```

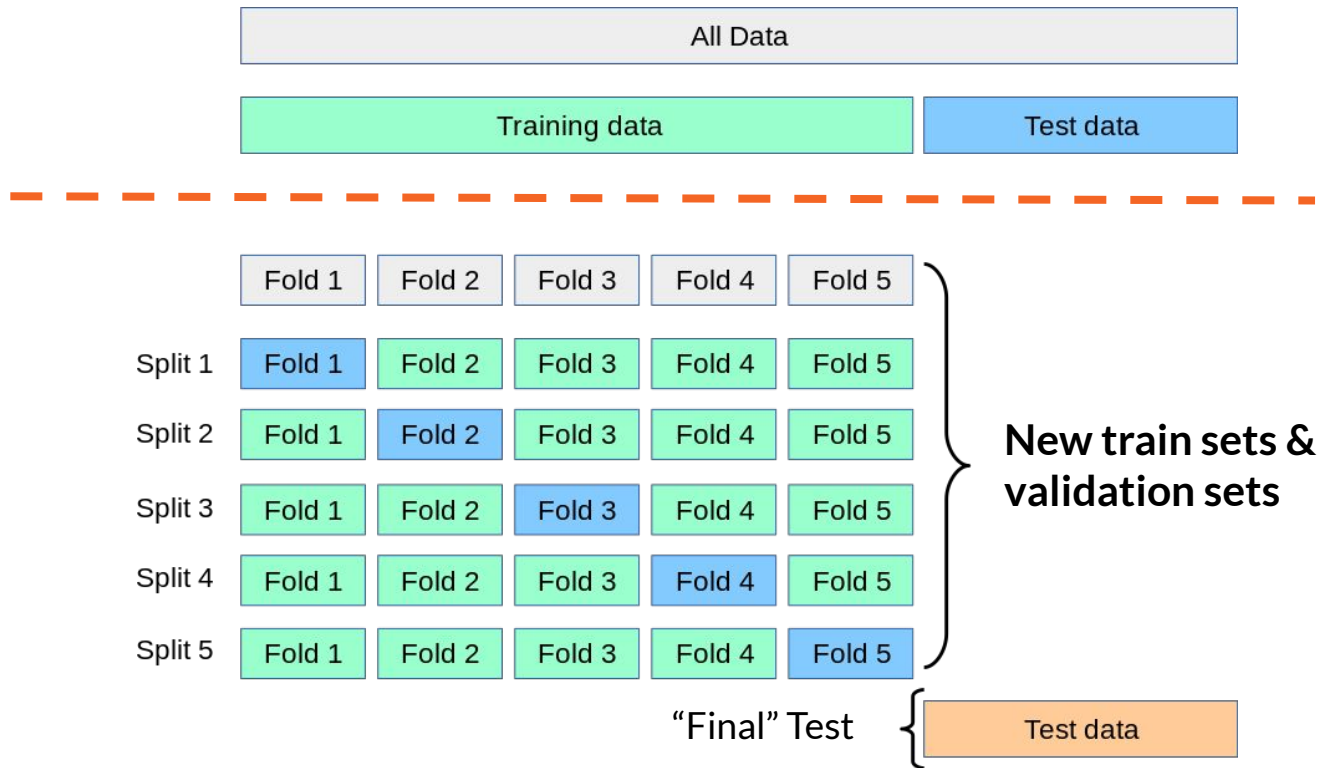
63/7/30 train/val/test split in second step (since the 70% train set gets split into 70*90% → train and 70*10% → validation)



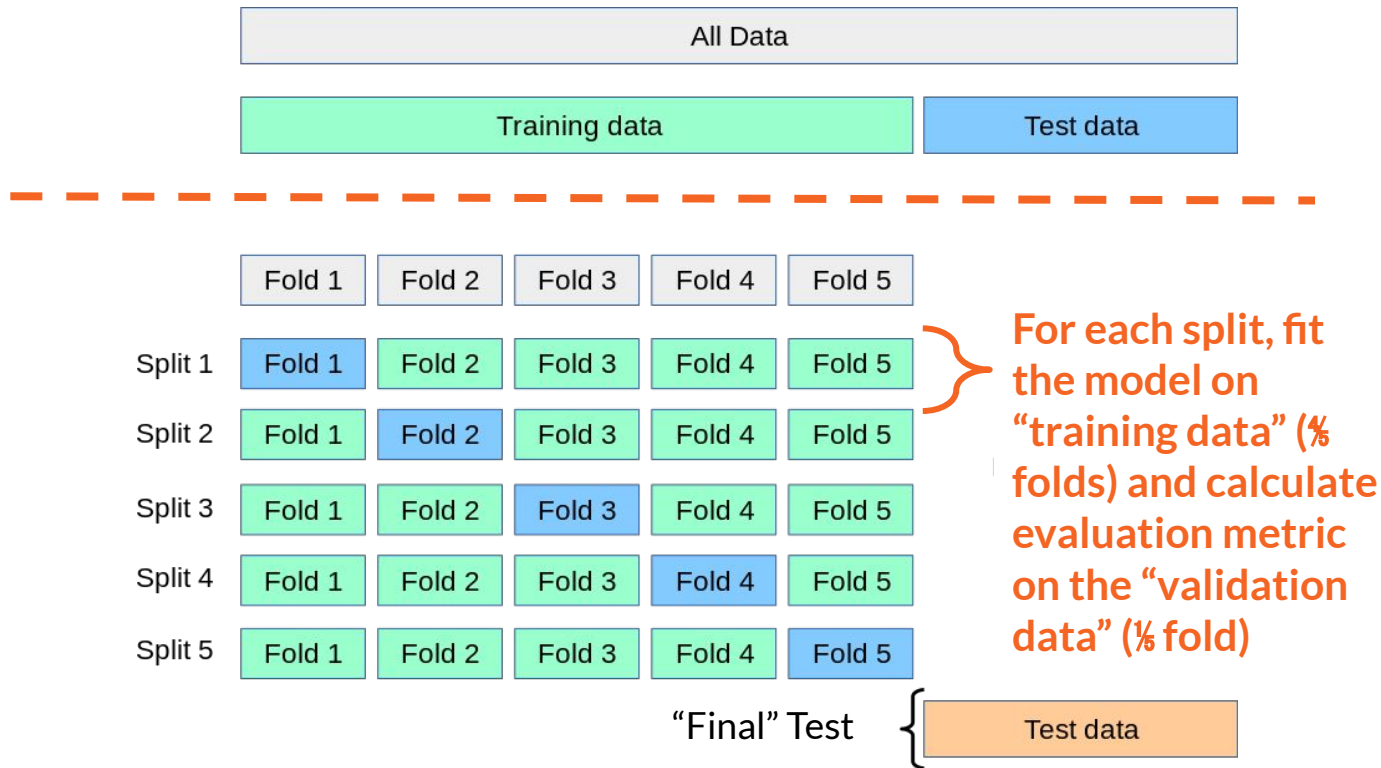
One Validation Set



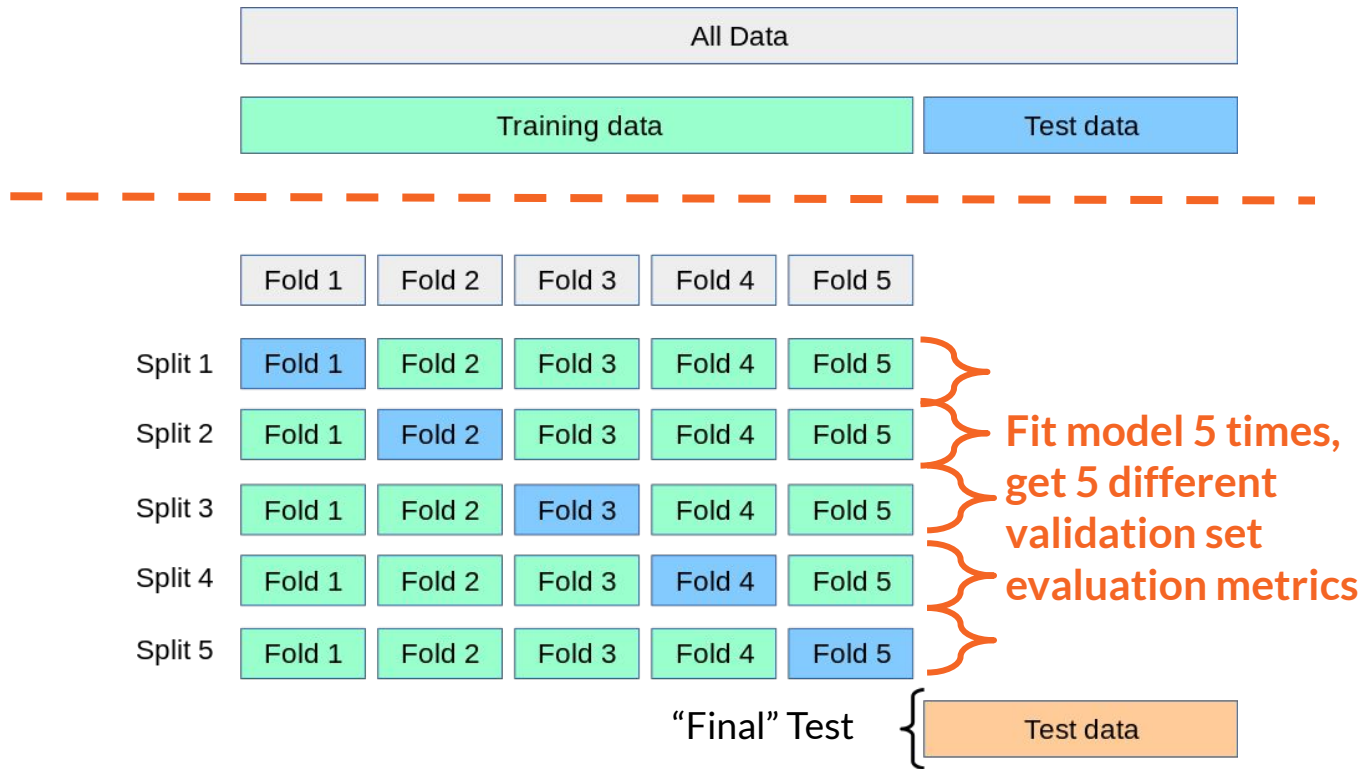
Cross Validation



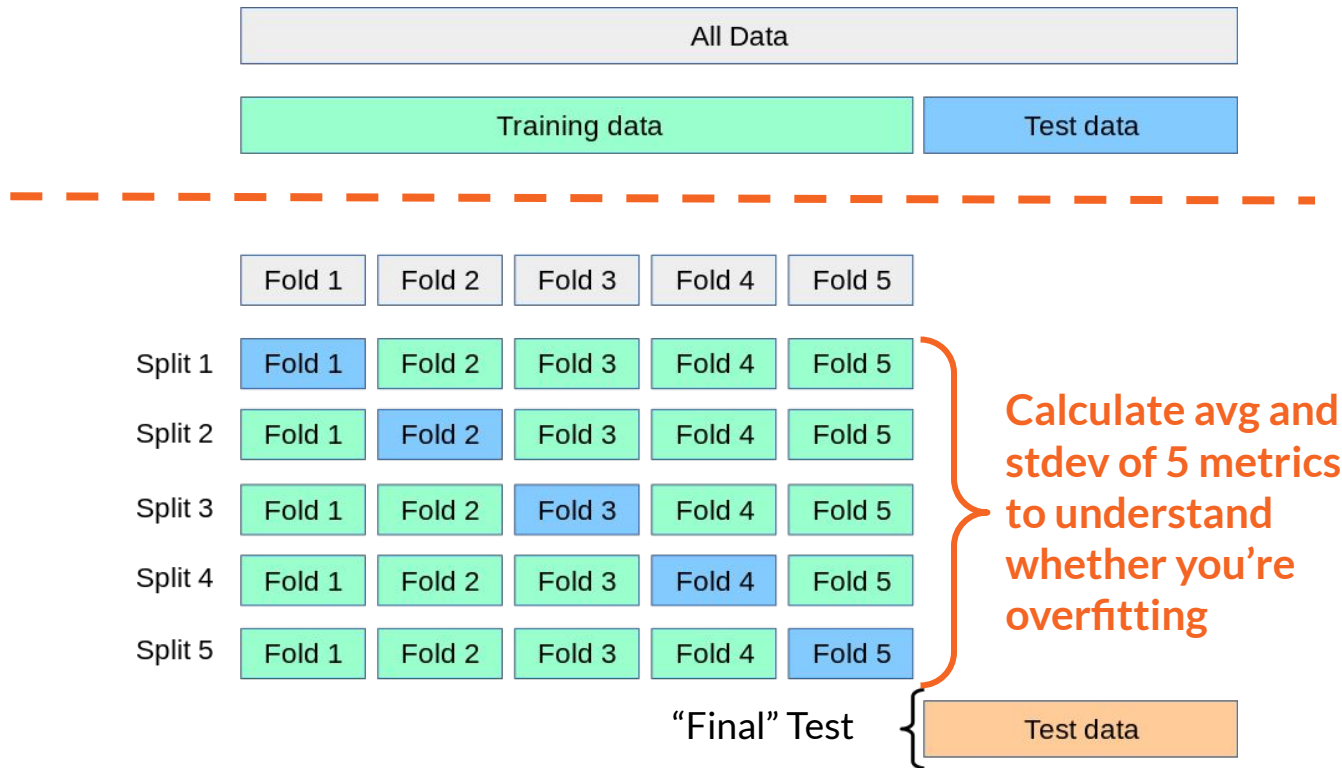
Cross Validation



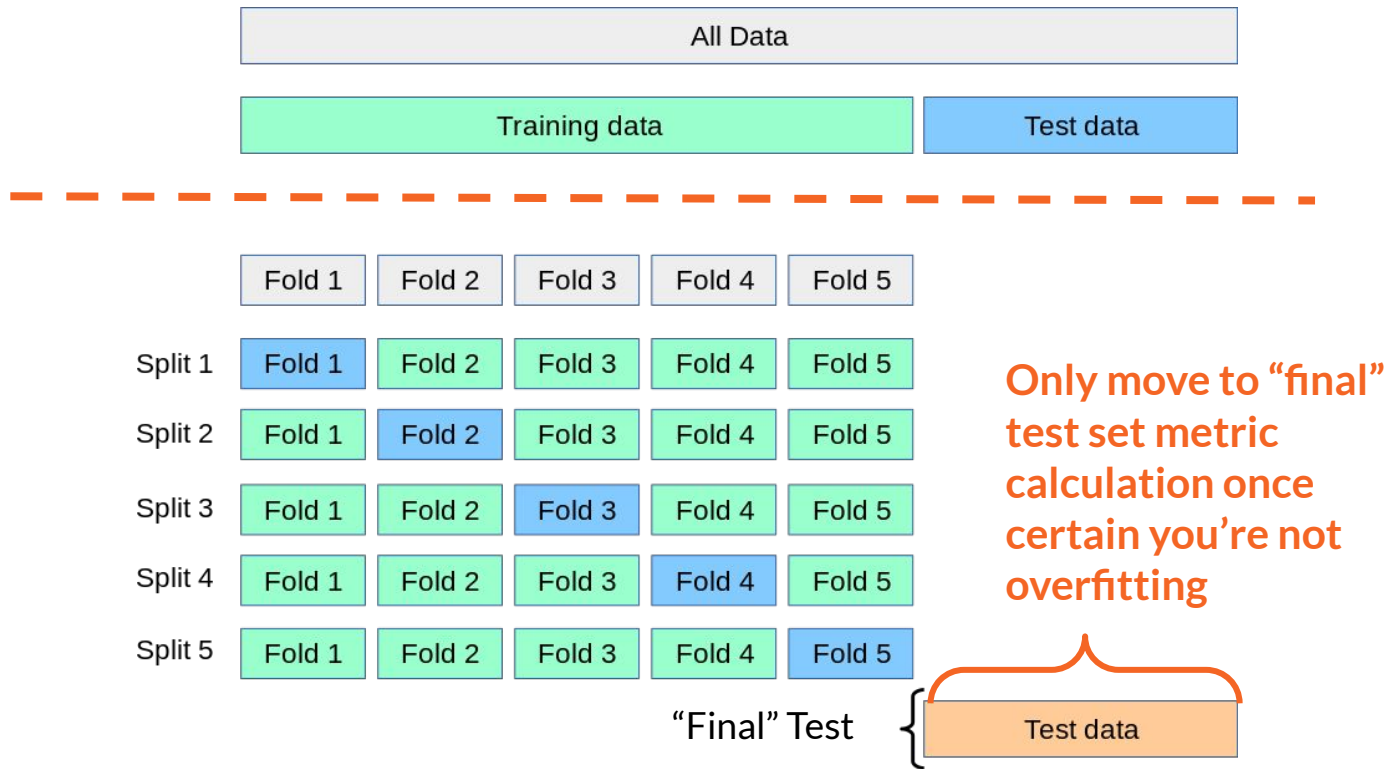
Cross Validation



Cross Validation



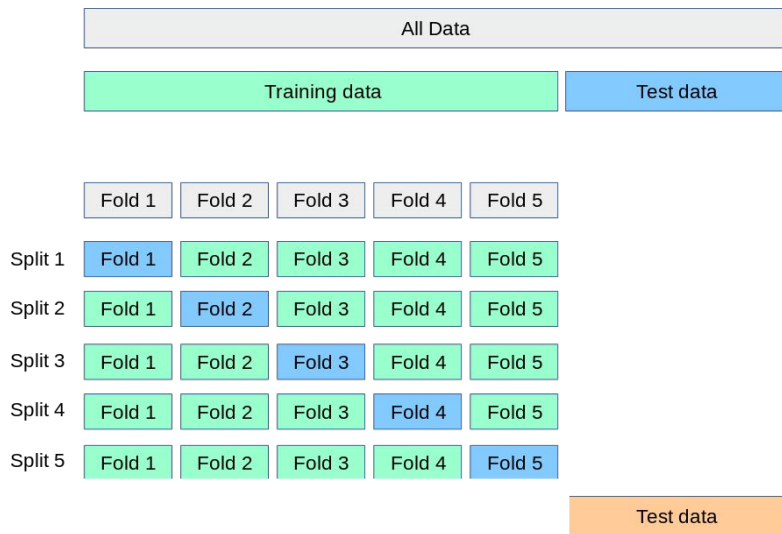
Cross Validation



Why CV?

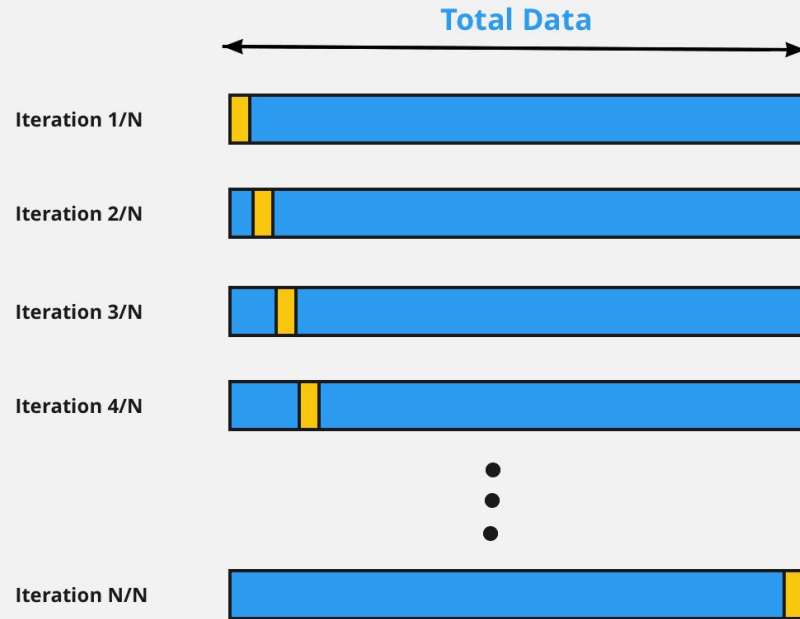
- Avoids pitfalls discussed earlier (getting lucky with one test set)
- Allows you to update your model without overfitting to a specific set of data
- Using CV with 5 folds → 5 values of an evaluation metric. Use summary stats to understand model performance **across different data samples**

Options for cross-validation

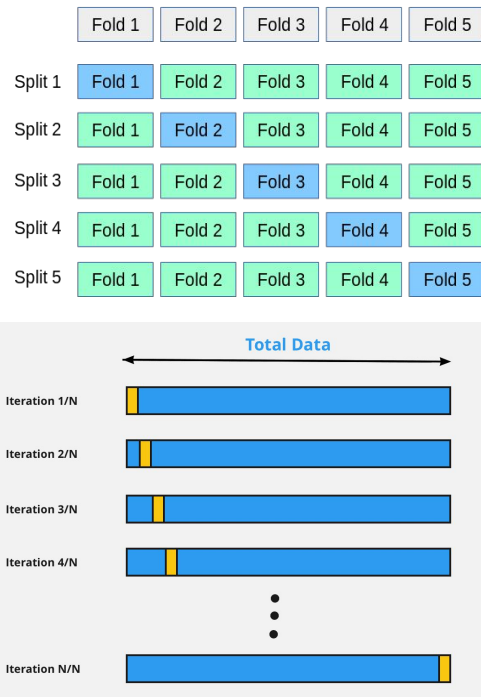


- Partitions or independent random splits
- Number of splits
- Special case: leave-one-out (LOO) cross-validation

LOOCV: Leave One Out Cross Validation

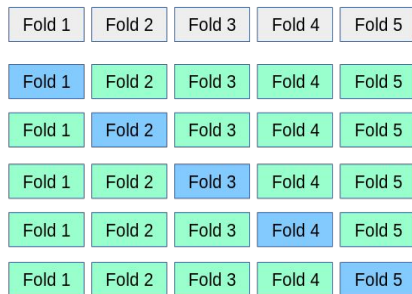


Cross Validation Variance

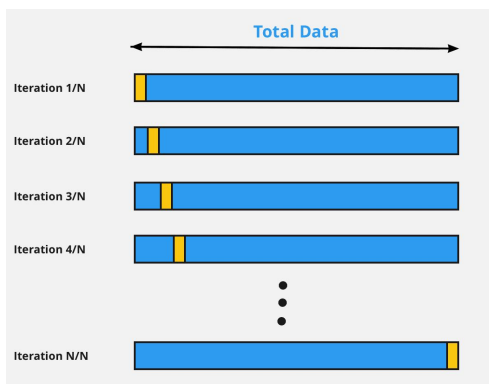


- If you increase the number of splits during cross-validation, does the *variance* of your estimated evaluation metric (the average of metrics across splits) tend to **increase or decrease**?

Cross Validation Variance



- If you increase the number of splits during cross-validation, does the *variance* of your estimated evaluation metric (the average of metrics across splits) increase or decrease?
- Generally, having more metrics to average (e.g. 1 per data point in the LOO case) → more likely to get outliers → **variance increases**



How to generate cross-validation splits

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test =
    train_test_split(X, y, test_size = 0.3)

for f in folds:
    X_train[f], X_valid[f], y_train[f], y_valid[f] =
        train_test_split(X_train, y_train, test_size = 0.1)
```

How to generate cross-validation splits

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3)  
for f in folds:  
    X_train[f], X_valid[f], y_train[f], y_valid[f] =  
        train_test_split(X_train, y_train, test_size = 0.1)
```

should I use a fixed random seed here?

How to generate cross-validation splits

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test =  
    train_test_split(X, y, test_size = 0.3)  
for f in folds:  
    X_train[f], X_valid[f], y_train[f], y_valid[f] =  
        train_test_split(X_train, y_train, test_size = 0.1)
```

No! That would just give me the same split folds times



How to generate cross-validation splits

```
from sklearn.model_selection import KFold
```

```
kf = KFold(n_splits=2)
```



New package

How to generate cross-validation splits

```
from sklearn.model_selection import KFold
```


```
kf = KFold(n_splits=2)
```

← Define how many
splits to use

How to generate cross-validation splits

```
from sklearn.model_selection import KFold  
  
X = ["a", "b", "c", "d"]  
kf = KFold(n_splits=2)
```

Defining some
sample data (think
of this as a column
of input data in the
training set)



How to generate cross-validation splits

```
from sklearn.model_selection import KFold
```

```
X = ["a", "b", "c", "d"]
```

```
kf = KFold(n_splits=2)
```

```
for train, test in kf.split(X):
```

```
    print(f"{train} {test}")
```

Split the input
data X



How to generate cross-validation splits

```
from sklearn.model_selection import KFold

X = ["a", "b", "c", "d"]

kf = KFold(n_splits=2)

for train, test in kf.split(X):

    print(f"{train} {test}")
```



Returns the
indices assigned
to train/test

How to generate cross-validation splits

```
from sklearn.model_selection import KFold

X = ["a", "b", "c", "d"]

kf = KFold(n_splits=2)

for train, test in kf.split(X):

    print(f"{train} {test}")
```

[2 3]	[0 1]
[0 1]	[2 3]

How to generate cross-validation splits

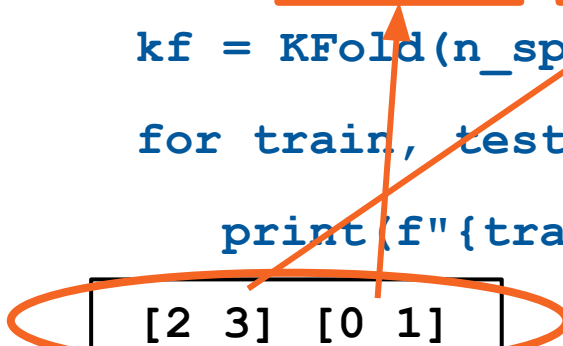
```
from sklearn.model_selection import KFold
```

```
X = ["a", "b", "c", "d"]
```

```
kf = KFold(n_splits=2)
```

```
for train, test in kf.split(X):
```

```
    print(f"{train} {test}")
```



[2 3]	[0 1]
[0 1]	[2 3]

Split 1: Train set indices 2,3 ("c", "d");
test set indices 0,1 ("a", "b")

How to generate cross-validation splits

```
from sklearn.model_selection import KFold  
  
X = ["a", "b", "c", "d"]  
kf = KFold(n_splits=2)  
for train, test in kf.split(X):  
    print(f"{train} {test}")
```

[2 3]	[0 1]
[0 1]	[2 3]

Split 2: Train set indices 0,1 ("a", "b");
test set indices 2,3 ("c", "d")

How to generate cross-validation splits

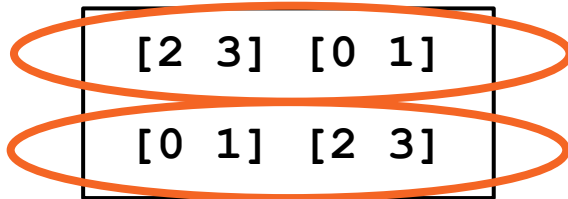
```
from sklearn.model_selection import KFold

X = ["a", "b", "c", "d"]

kf = KFold(n_splits=2)

for train, test in kf.split(X):

    print(f"{train} {test}")
```



[2 3]	[0 1]
[0 1]	[2 3]

Calculate evaluation metrics on split 1

Calculate evaluation metrics on split 2

How to generate cross-validation splits

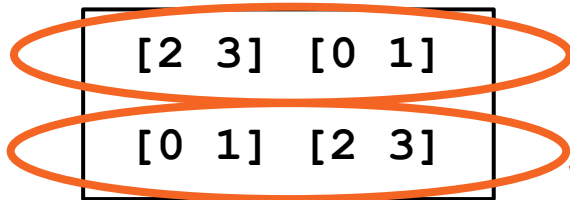
```
from sklearn.model_selection import KFold

X = ["a", "b", "c", "d"]

kf = KFold(n_splits=2)

for train, test in kf.split(X):

    print(f"{train} {test}")
```



Calculate the mean / stdev of
evaluation metrics across splits

What could go wrong with train/test?

Solution:
Cross-Validation

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)
- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

When you realize k-fold cross validation can only validate your models, not yourself...



What could go wrong with train/test?

- Train set proportion too big (test set not broad enough)
- Test set proportion too big (not enough training data)
- Test set data too similar/dissimilar to train set data (overestimating/underestimating generalization)
- Having ordered rows in a df (shouldn't randomize x-axis order of time series)

Solution:
Sliding Windows

Sliding Windows

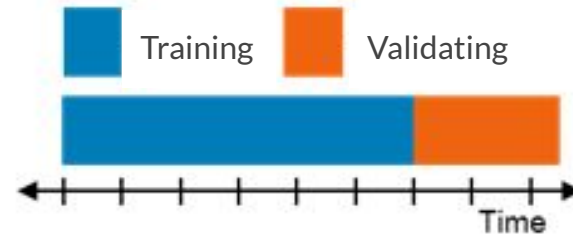
- For time series: take the concept of cross validation, but tweak it so that you can't use the **future** to model the **past**
 - This is cheating (“data leakage”) and will give you overly-good results!

Sliding Windows

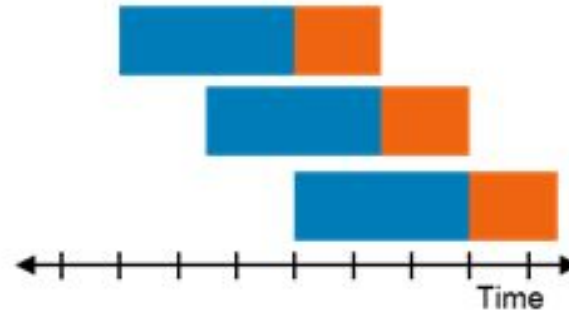
- For time series: take the concept of cross validation, but tweak it so that you can't use the **future** to model the **past**
 - This is cheating (“data leakage”) and will give you overly-good results!
- Instead, **slide** your test/validation sets forward by a “time **window**” unit

Sliding Windows

Time-based Estimation

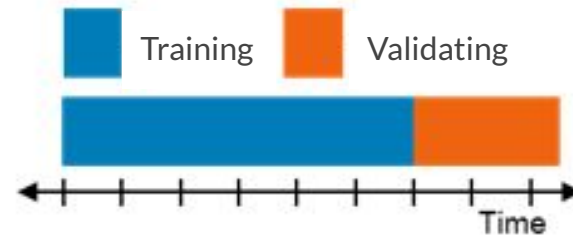


Time-based cross-validation



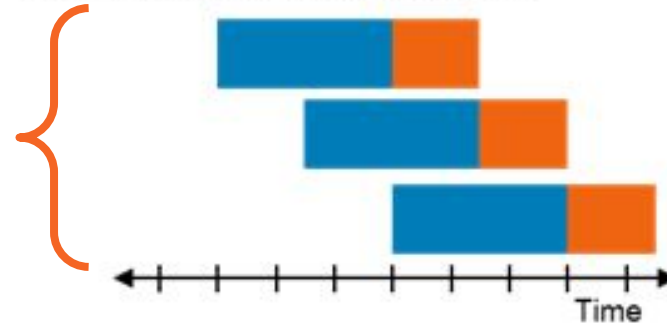
Sliding Windows

Time-based Estimation



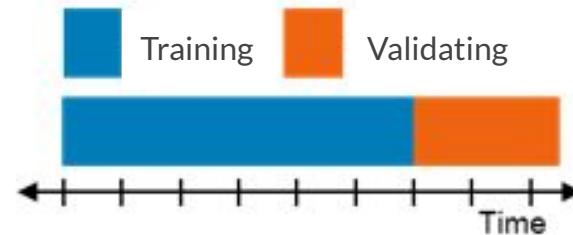
Time-based cross-validation

Multiple splits, each one with a different train and validation set



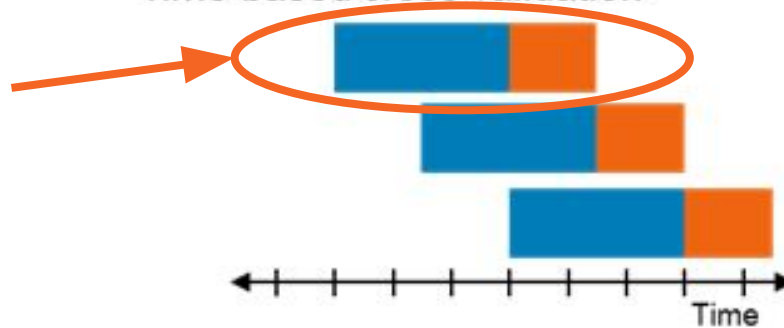
Sliding Windows

Time-based Estimation



Time-based cross-validation

Never peeking at “future”
test data when training
the model



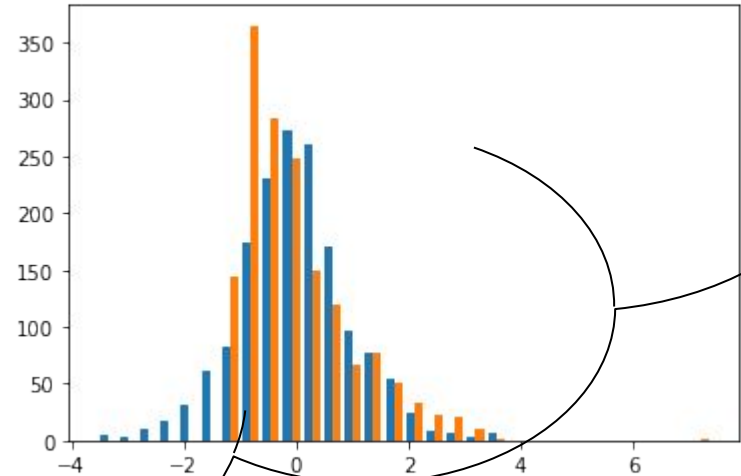
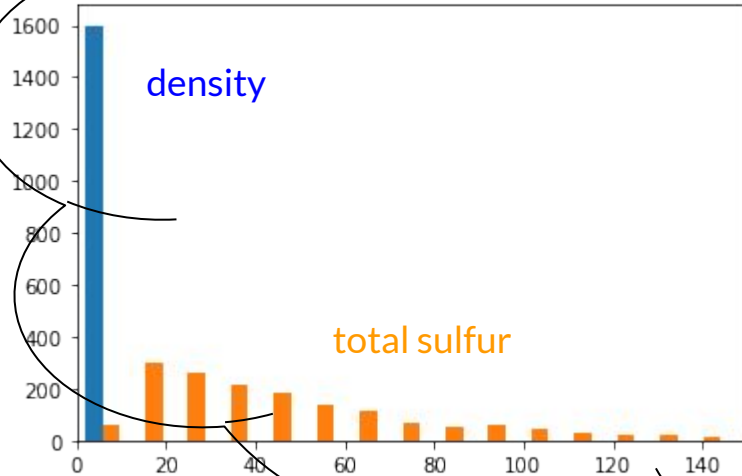
Data Leakage

- We don't want to “leak” information between the train / validation / test sets, in order to keep the val / test set evaluations **fully separate** from the train set evaluations
 - “Leaking” via time

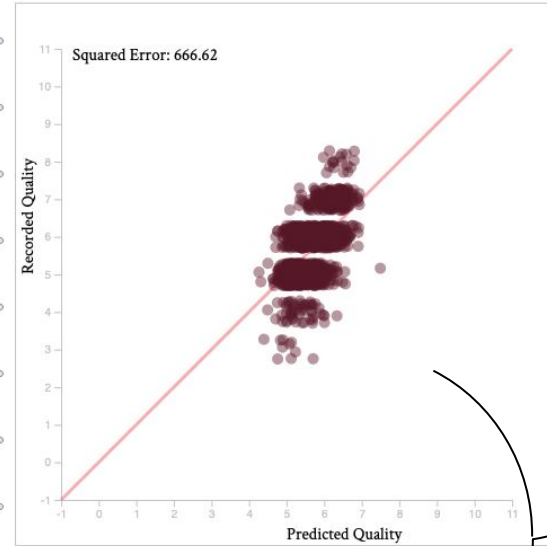
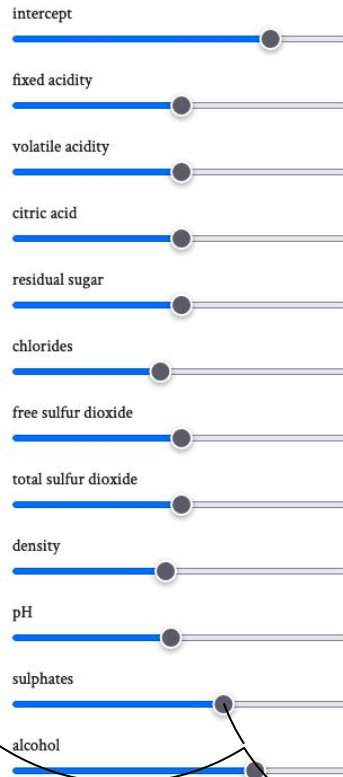
Data Leakage

- We don't want to “leak” information between the train / validation / test sets, in order to keep the val / test set evaluations **fully separate** from the train set evaluations
 - “Leaking” via time
 - “Leaking” via normalization

Input variables are on different scales: *normalize with z-scores*



RESET **NORMALIZE** HINT Learning Rate: 0.01 STEP



Normalizing → data leakage?

- Which of the following would prevent data leakage when normalizing data?

A: First generate train/test splits, then normalize within train set, and separately normalize within test set

B: First normalize all the data, then generate train/test splits

Normalizing → data leakage?

- Which of the following would prevent data leakage when normalizing data?

A: First generate train/test splits, then normalize within train set, and separately normalize within test set

If you instead normalize across all the data, the mean and stdev you calculate will incorporate data across both train/test sets, effectively “leaking” information

How to do this CV stuff in Python?

Check the Scikit Learn documentation:
https://scikit-learn.org/stable/modules/cross_validation.html